Machine Learning on OpenShift

From Data Scientist to Application Developer

Michael McCune
Principal Software Engineer

Red Hat

Forecast

- Introduction
- Technology Review
- Building Intelligent Applications
- Lessons Learned
- Next Steps

Who is this guy?

- Joined Red Hat 6 years ago
- Full stack, from embedded to orchestration
- Emerging technology at Red Hat
- Big Data on OpenStack and OpenShift

We are talking about a journey

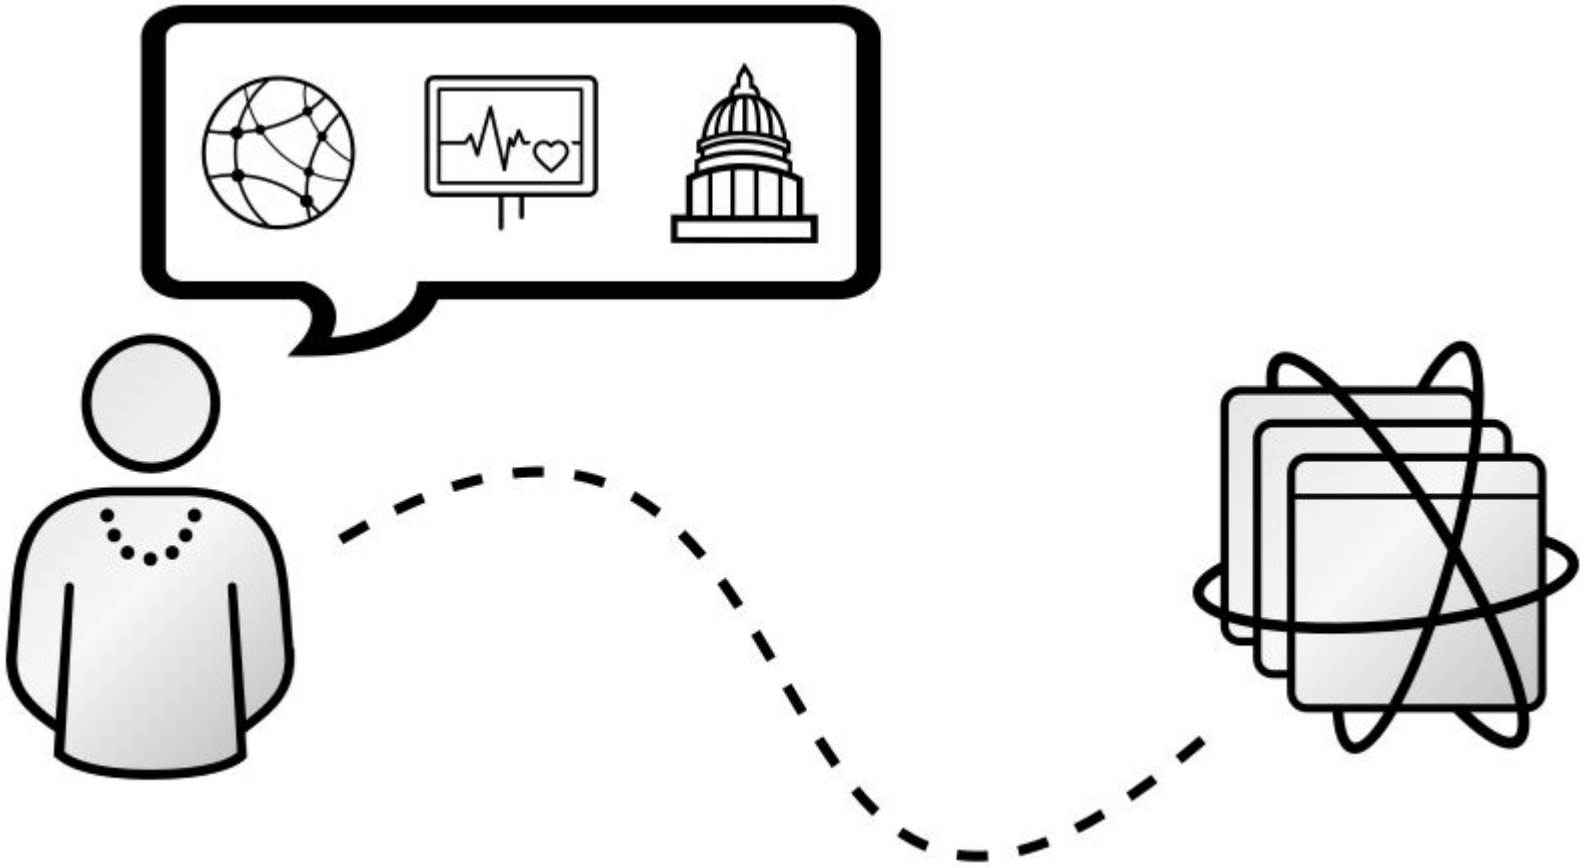

Red Hat

photography by Sam Hawley

Red Hat

photography by Diliff and Janke

Red Hat

Level Set

Red Hat

Machine Learning

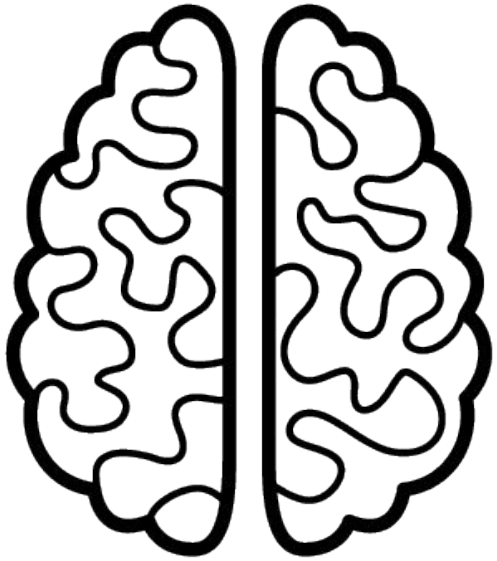

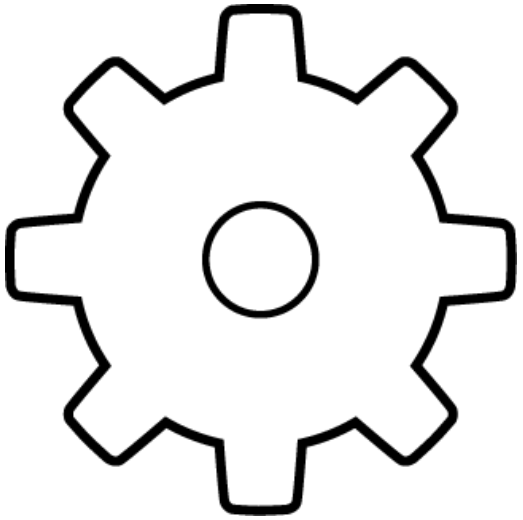

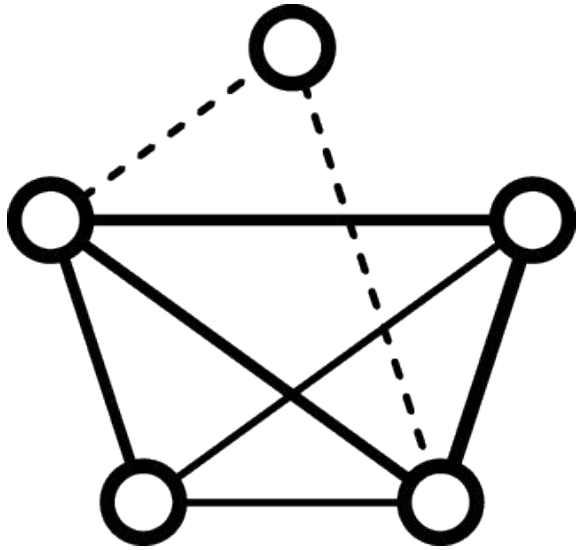

Identifying Trends

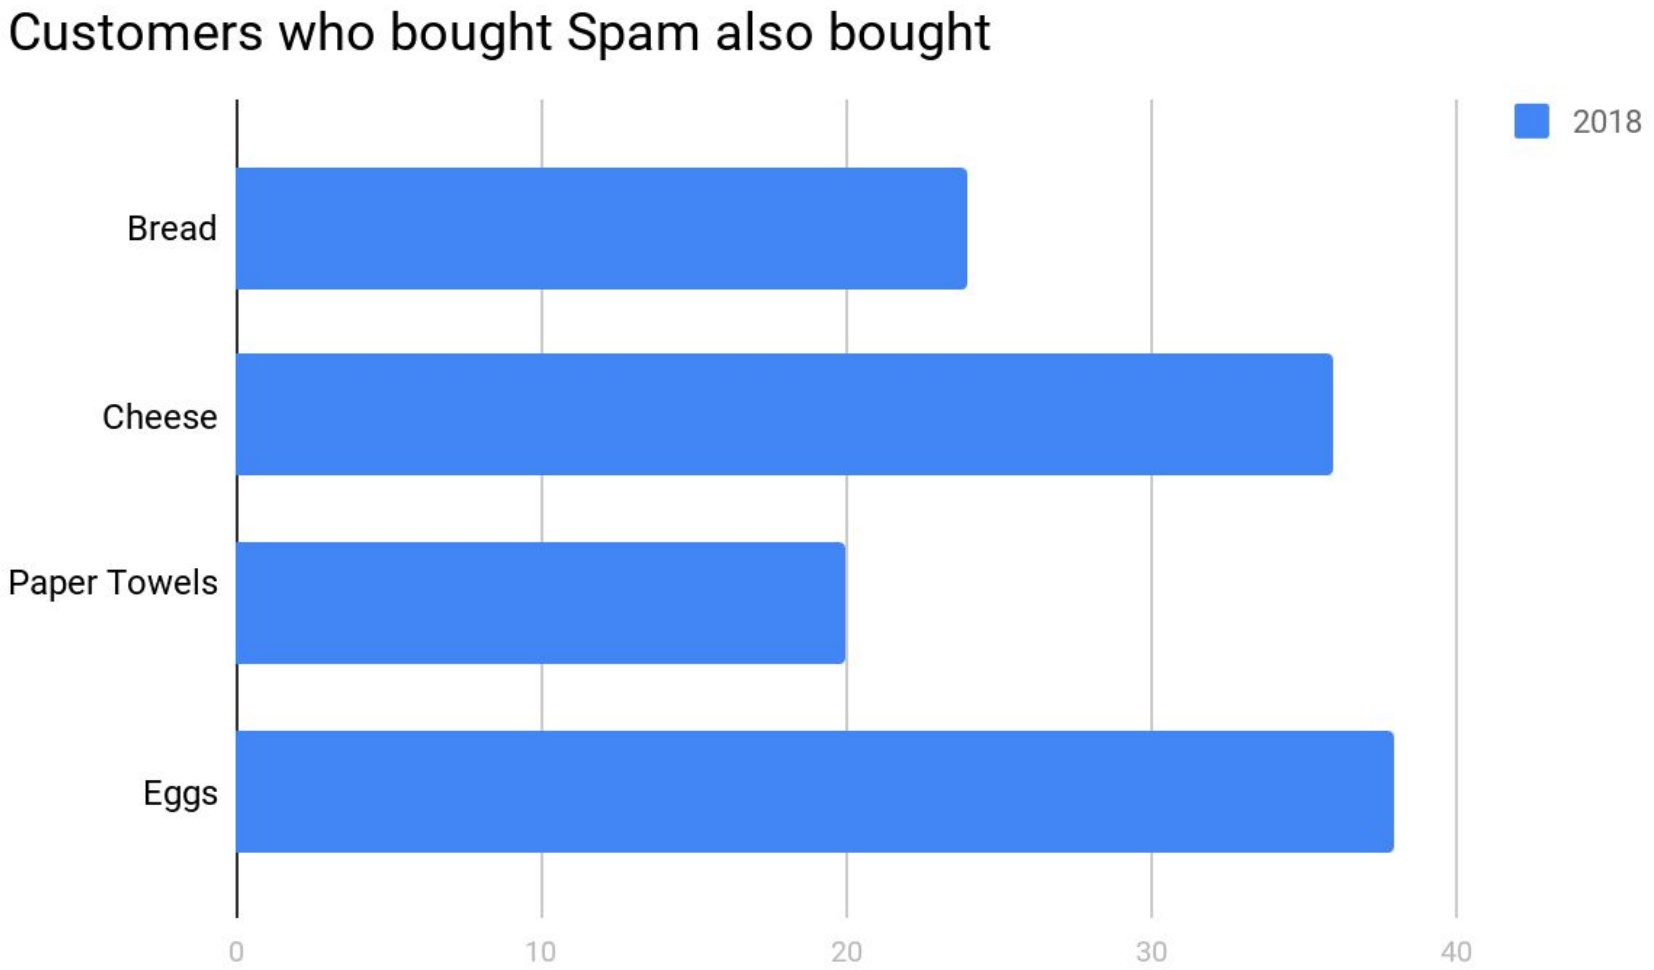

Red Hat

Predicting Results

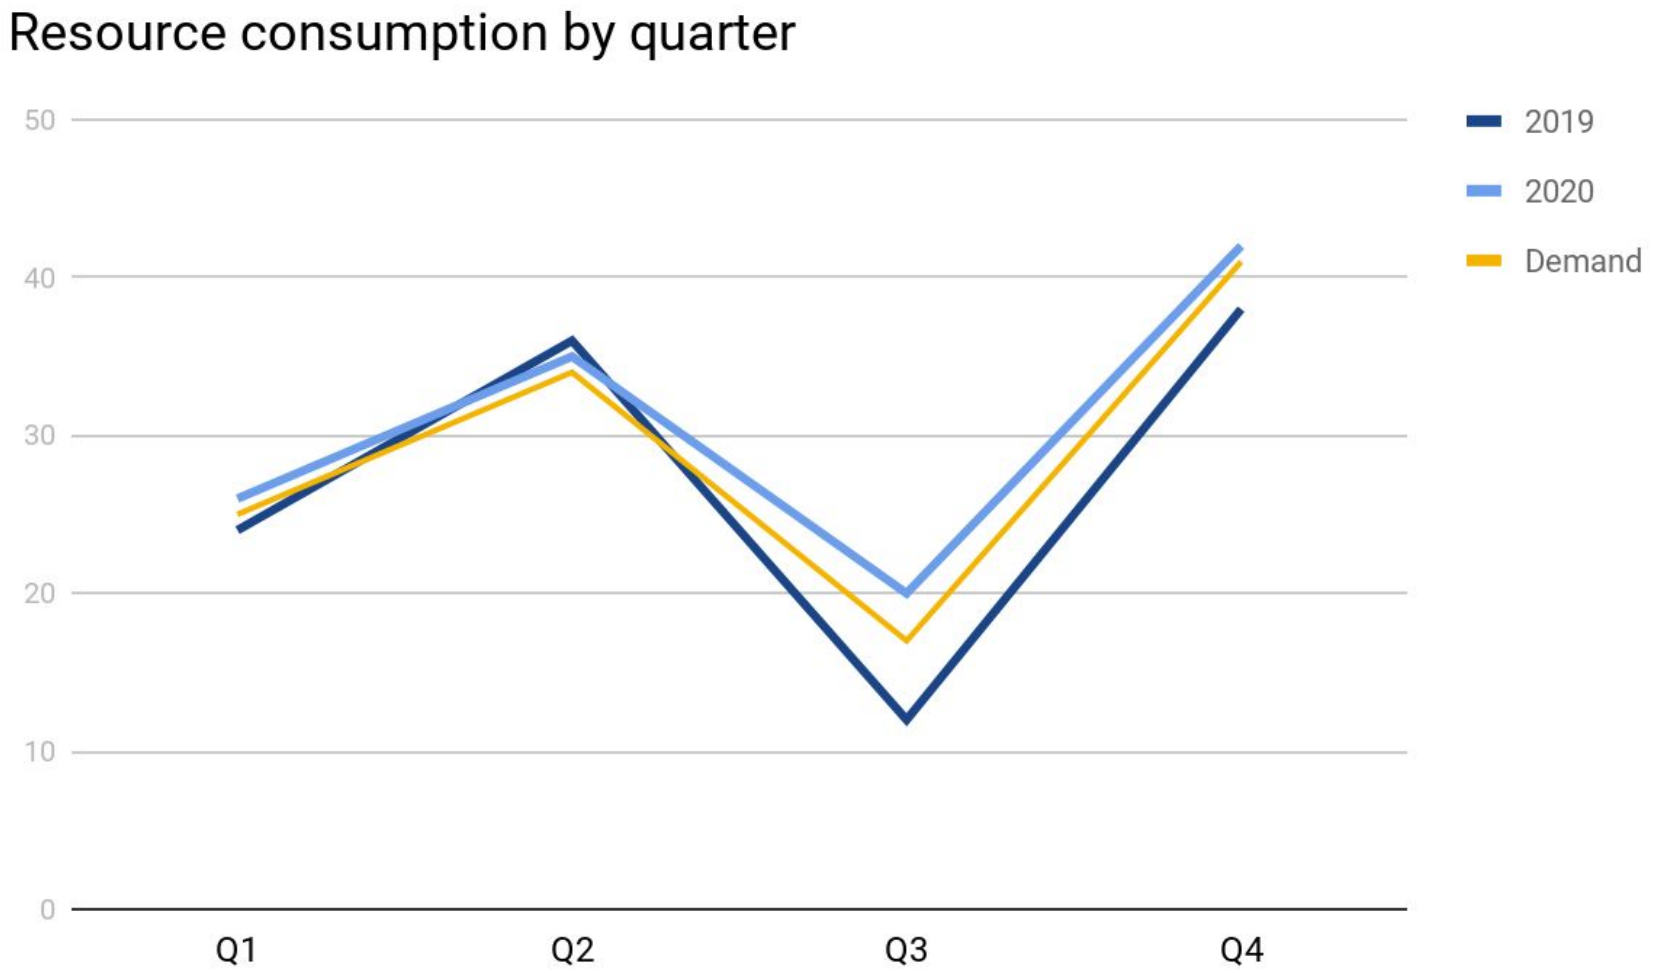

Red Hat

Estimating Values

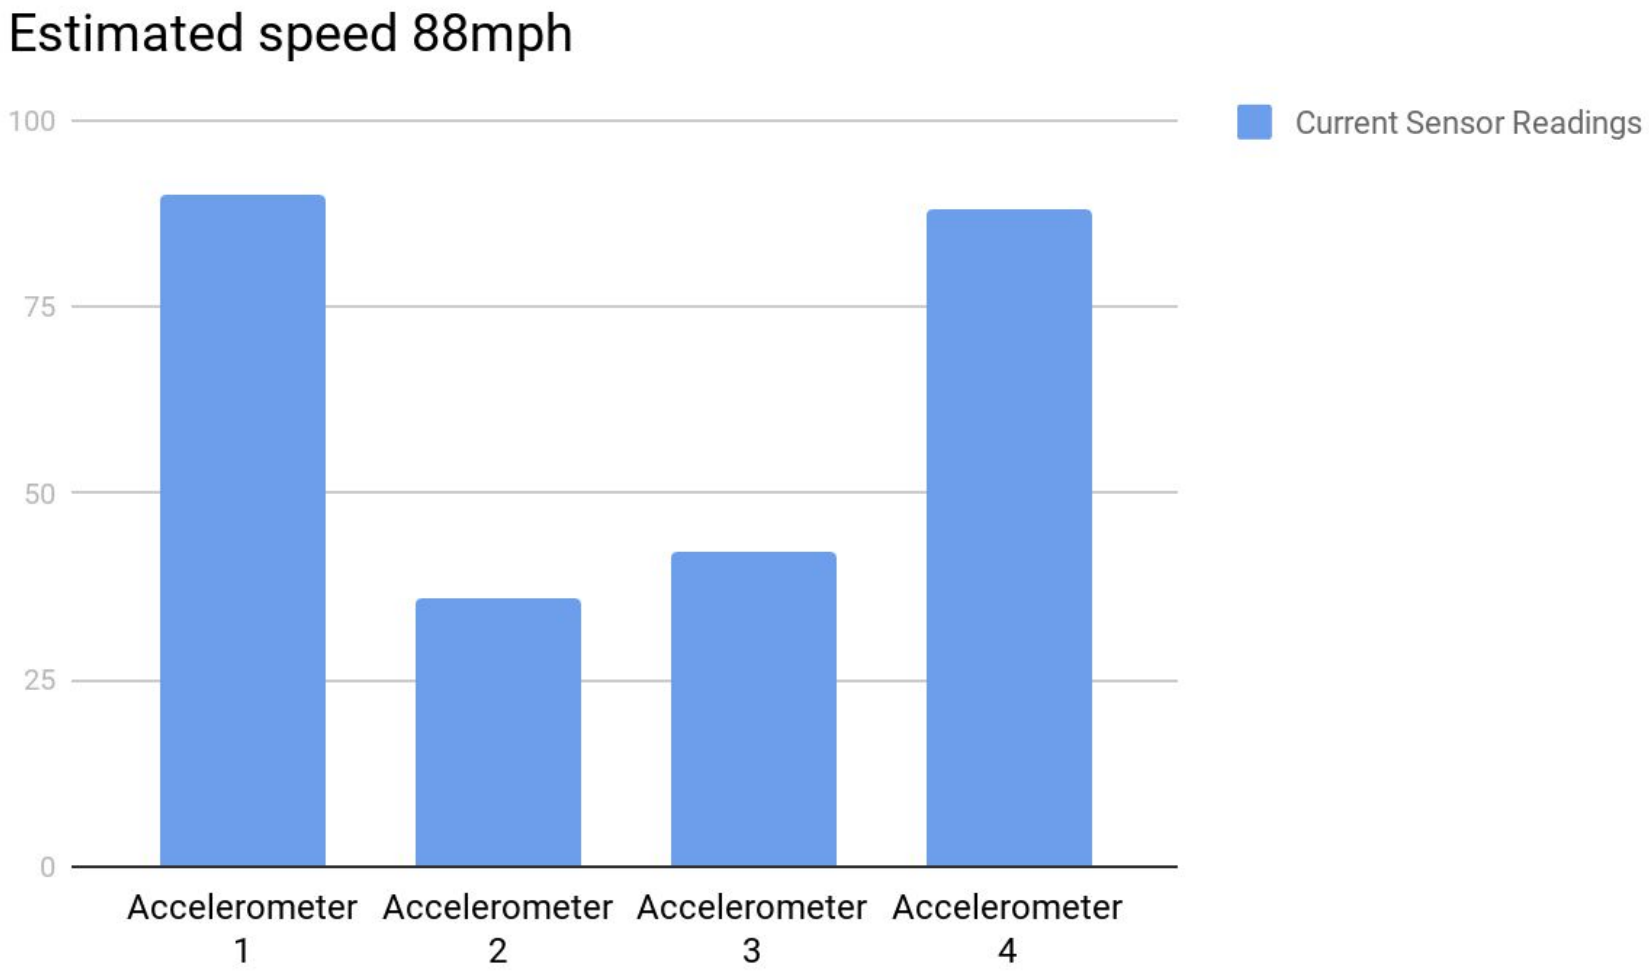

Red Hat

Kubernetes

Red Hat

Let’s talk containers

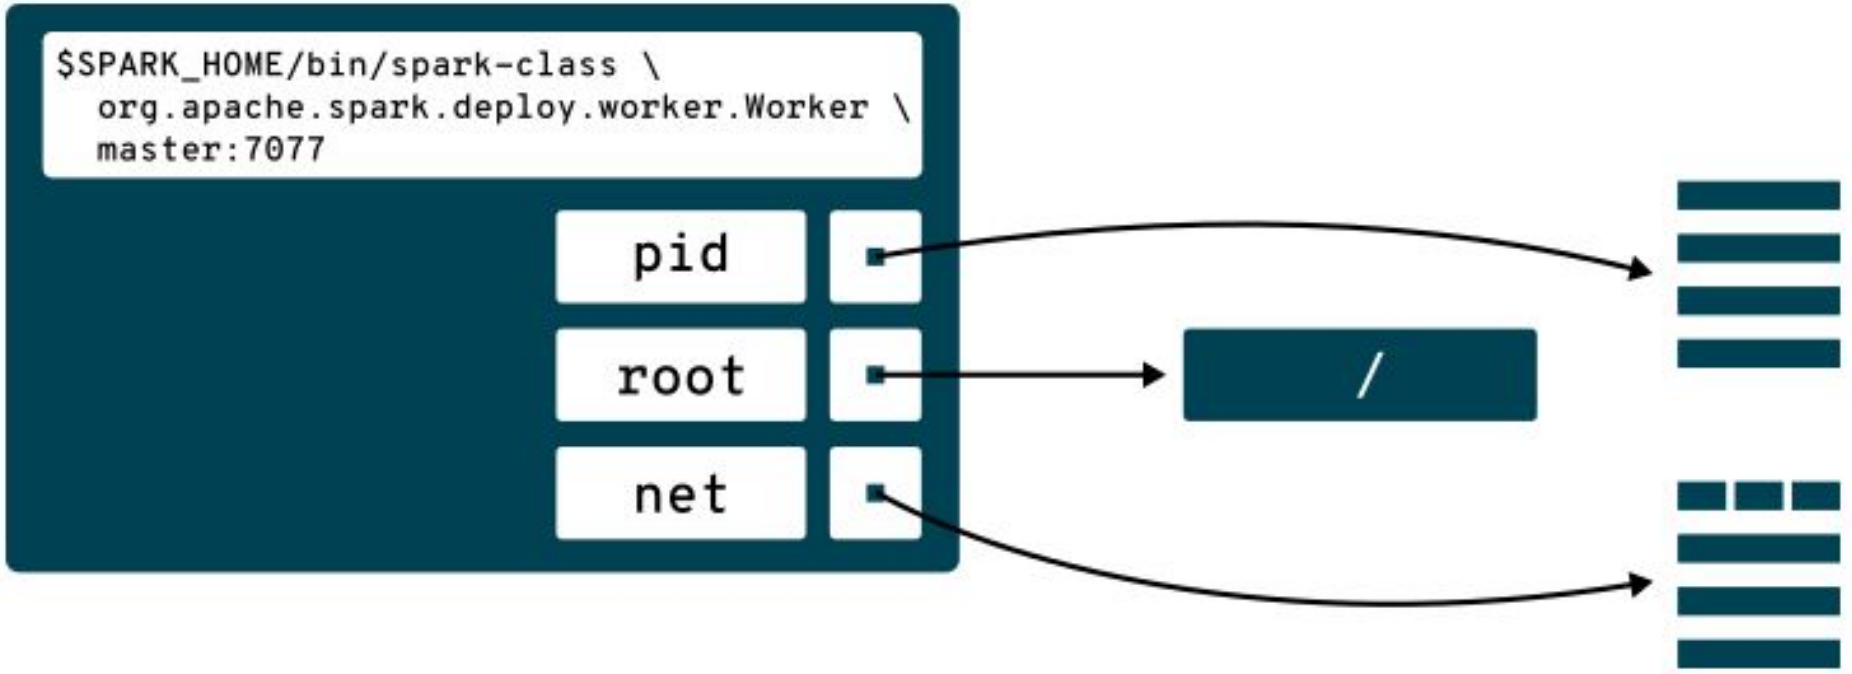

Red Hat

Let’s talk containers

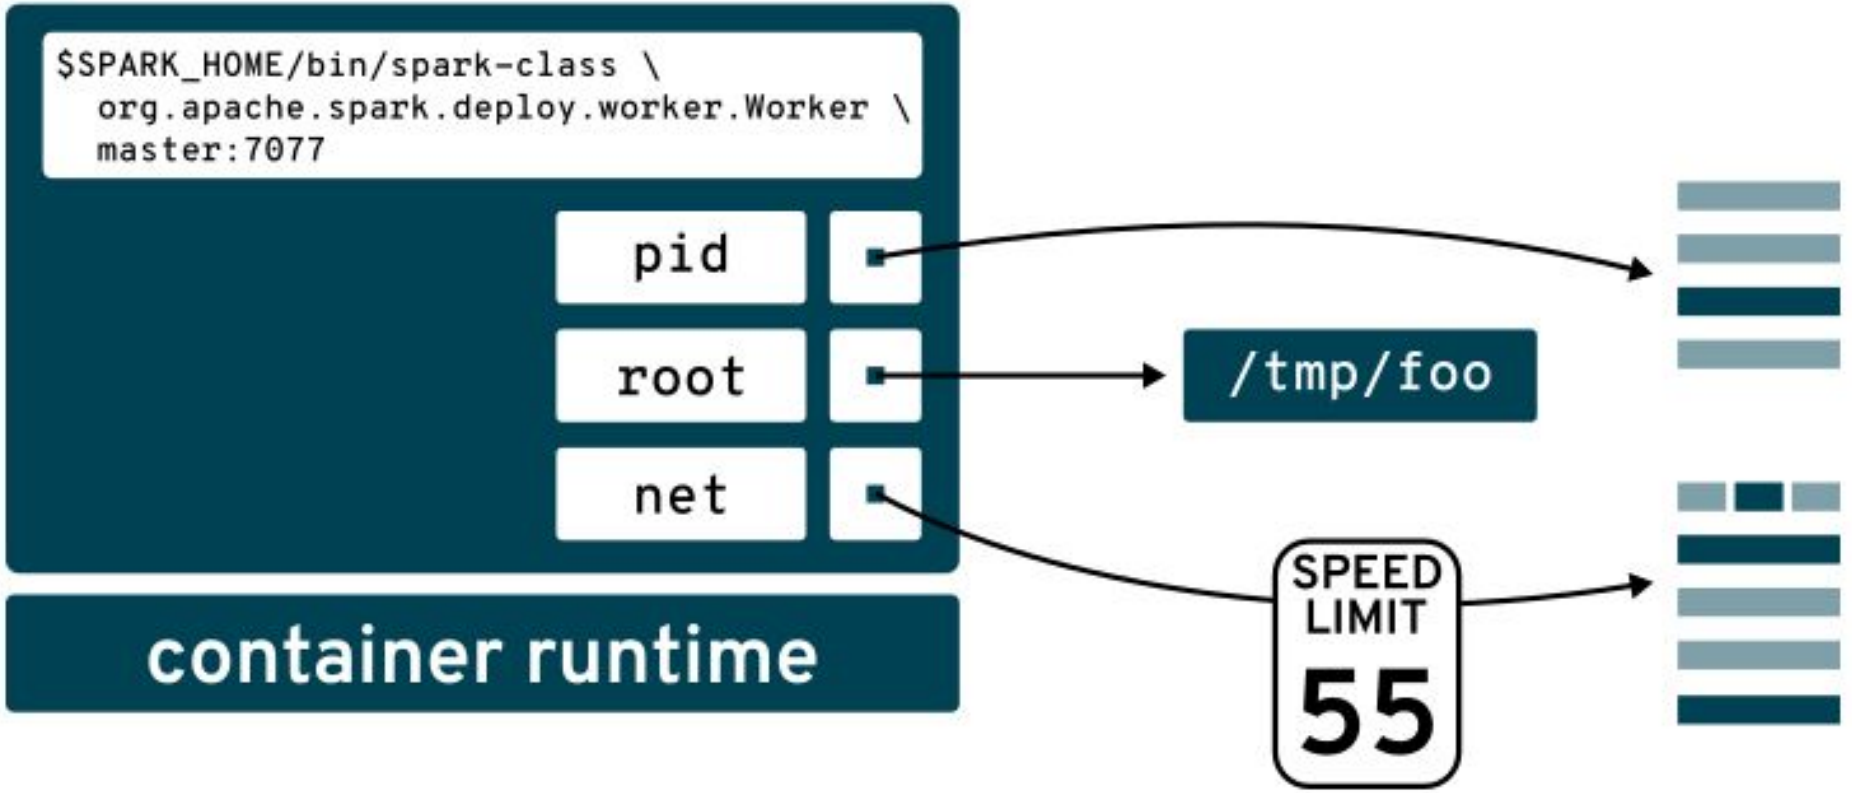

Red Hat

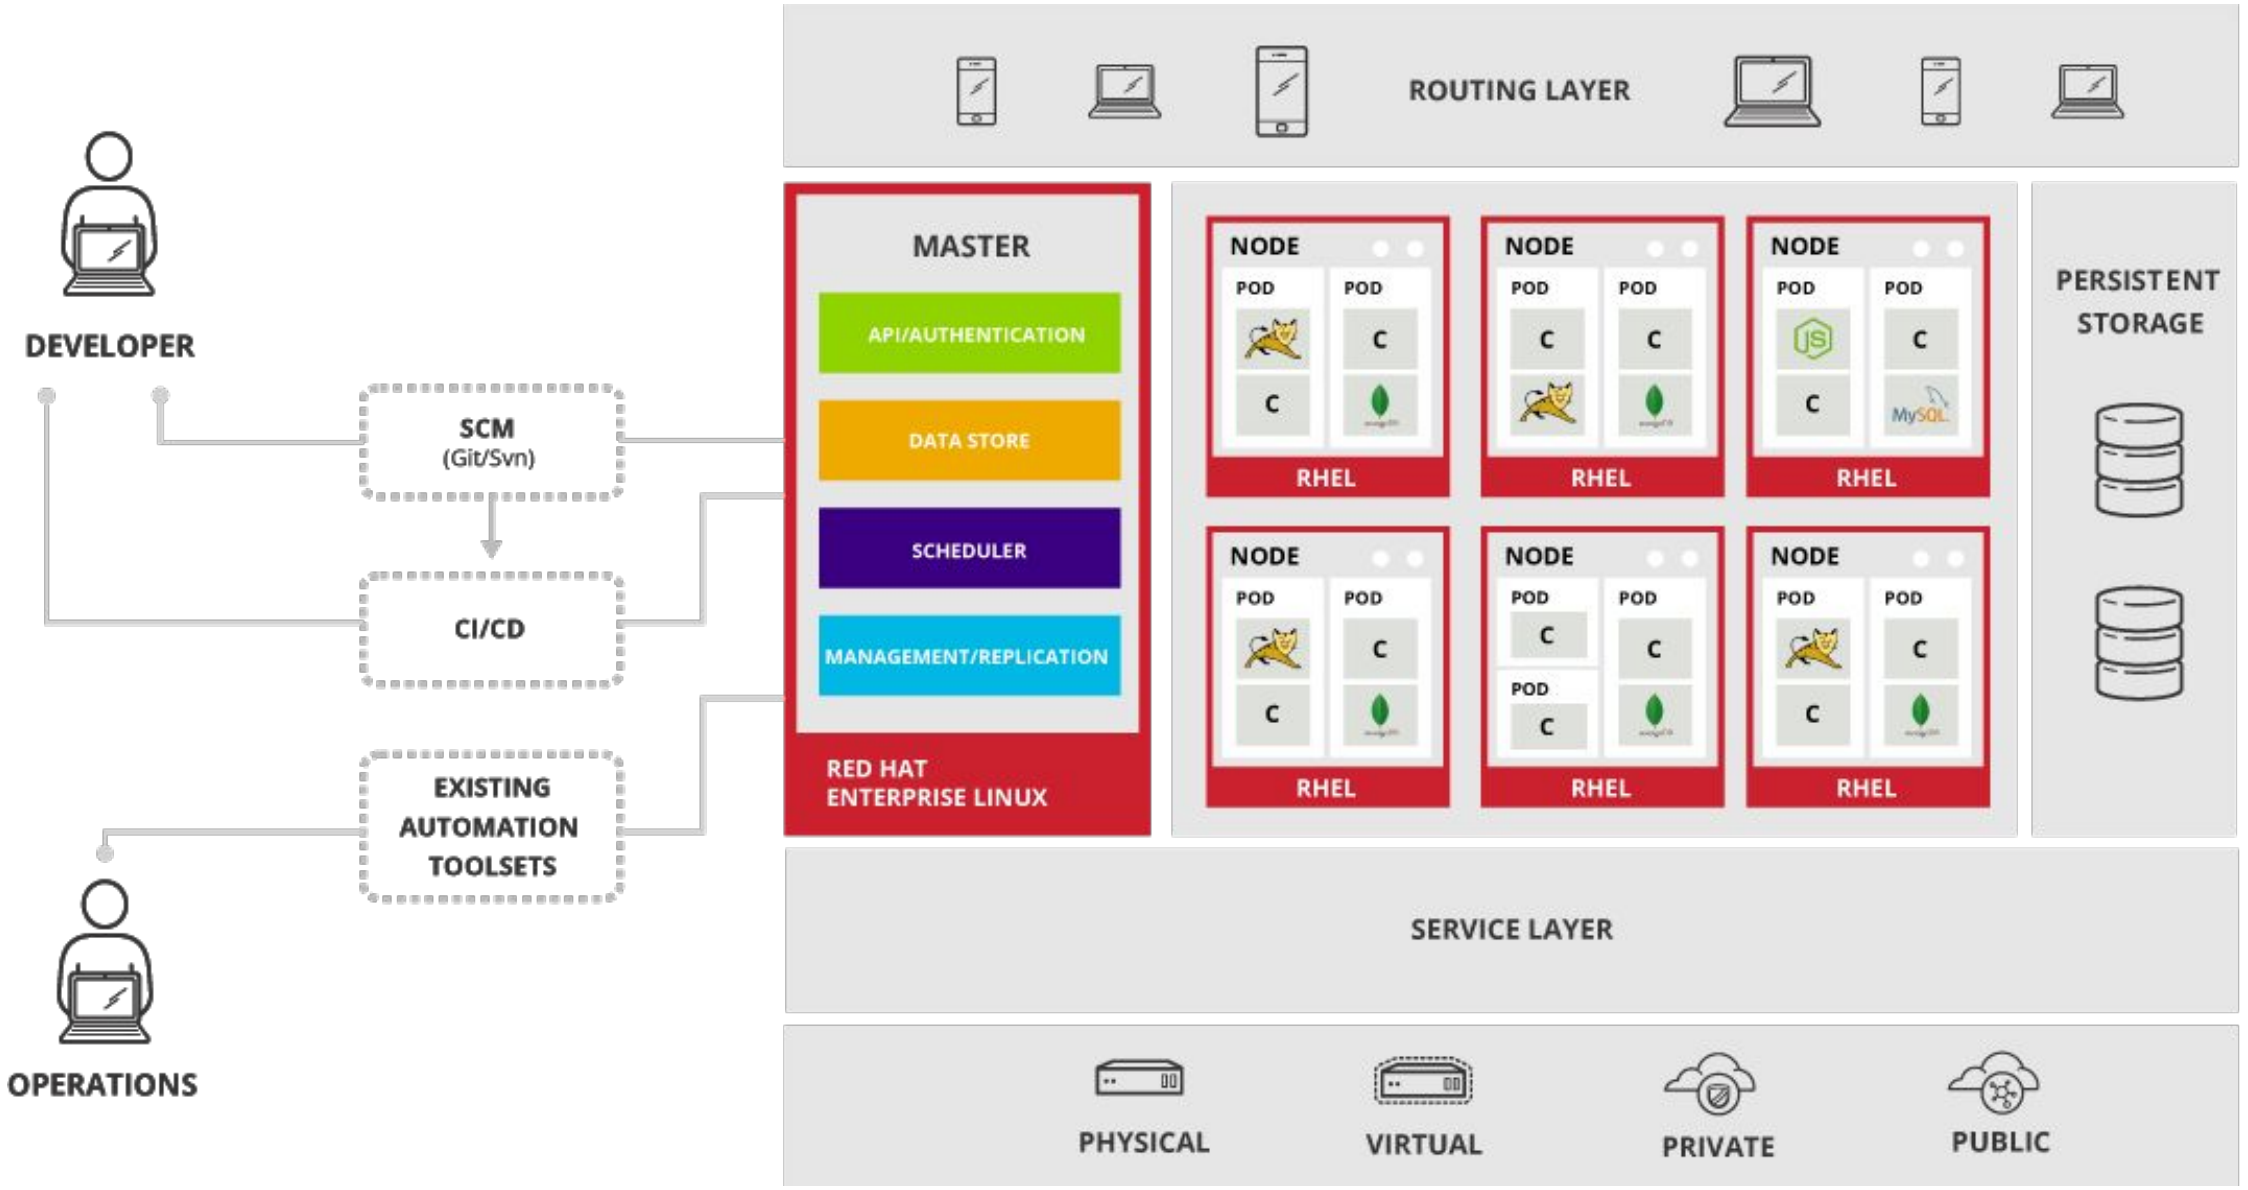
ROUTING LAYER
DEVELOPER
SCM
(Git/Svn)
CI/CD
EXISTING
AUTOMATION
TOOLSETS
OPERATIONS
MASTER
API/AUTHENTICATION
DATA STORE
SCHEDULER
MANAGEMENT/REPLICATION
RED HAT
ENTERPRISE LINUX
NODE
POD
C
RHEL
PERSISTENT
STORAGE
SERVICE LAYER
PHYSICAL
VIRTUAL
PRIVATE
PUBLIC

Red Hat

Microservice Architectures

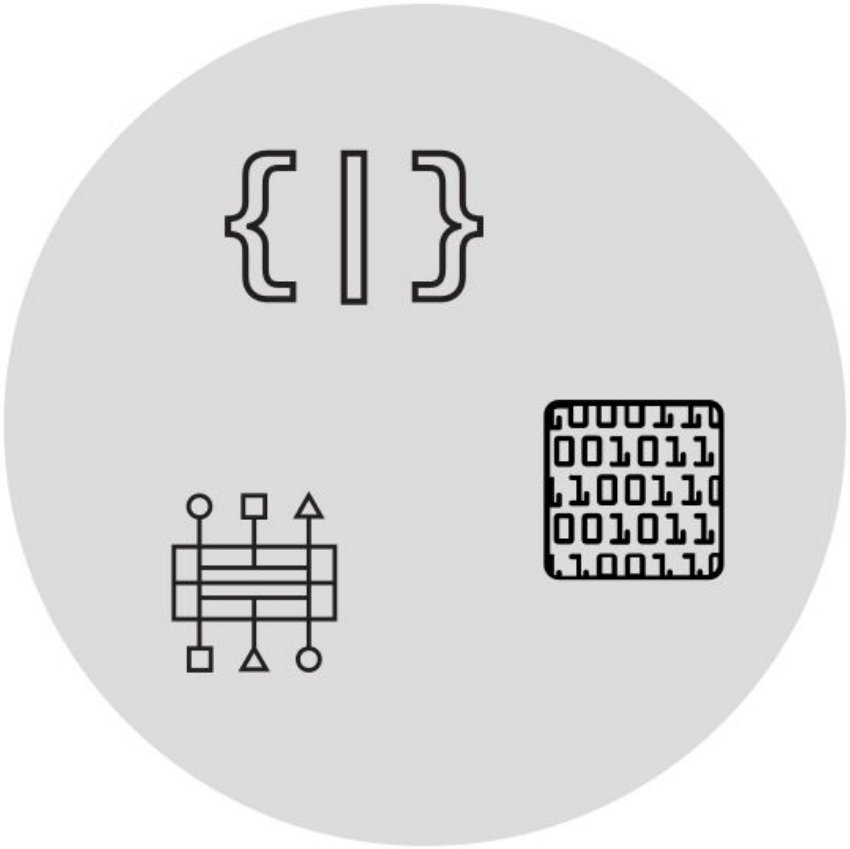

Red Hat

Composable

Scalable

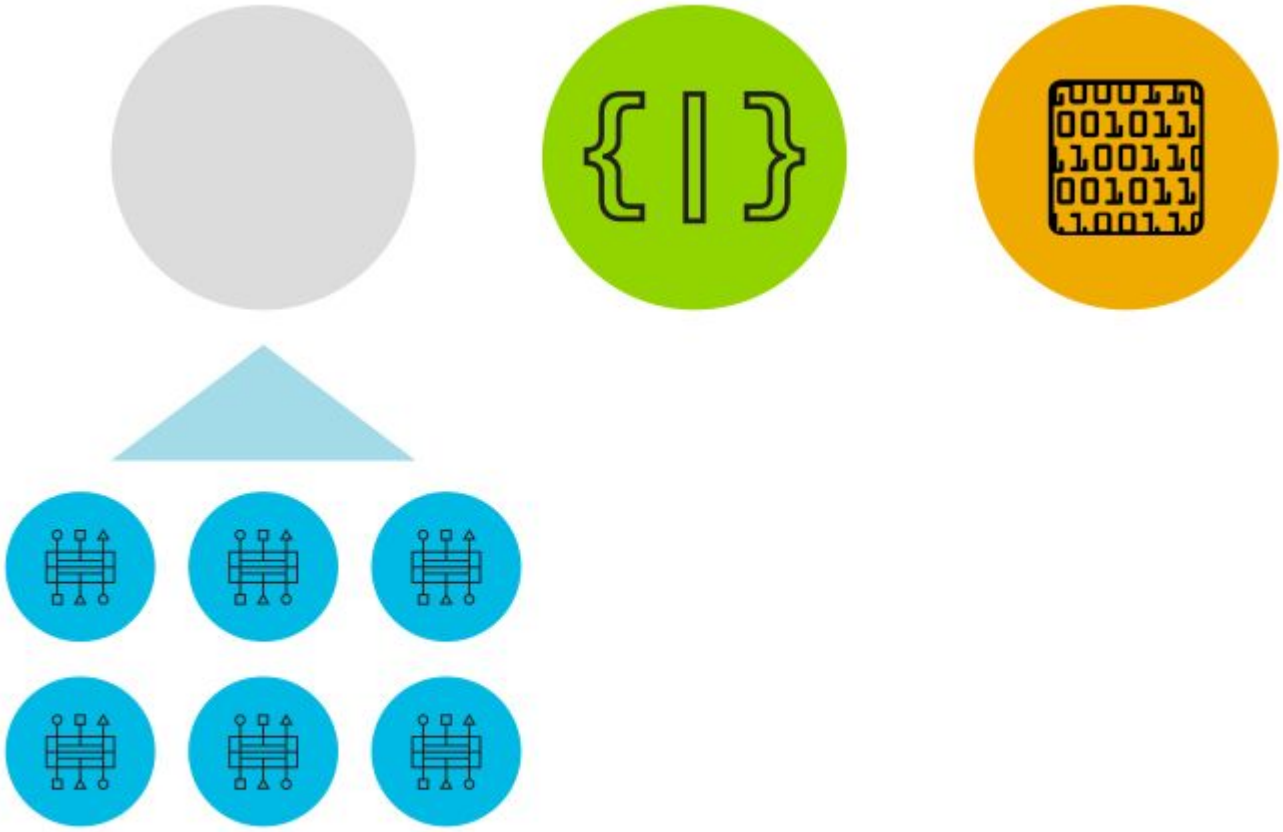

Flexible

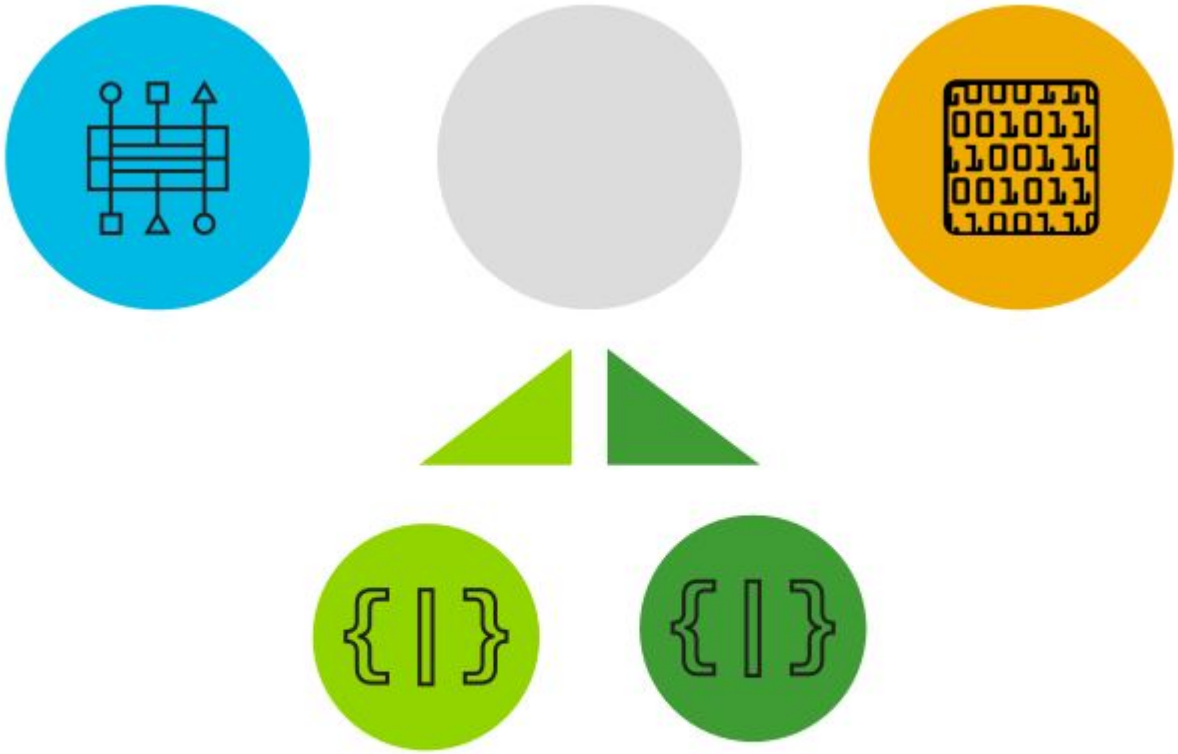

Resilient

Intelligent Applications

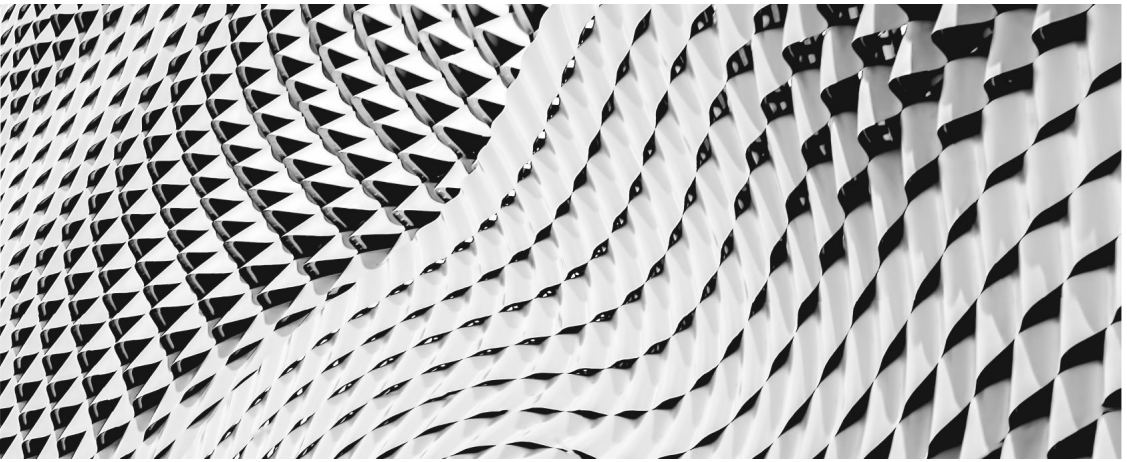

Intelligent applications **collect and learn from data** to provide **improved functionality** with **longevity and popularity**.

Red Hat

Intelligent Application Pipeline

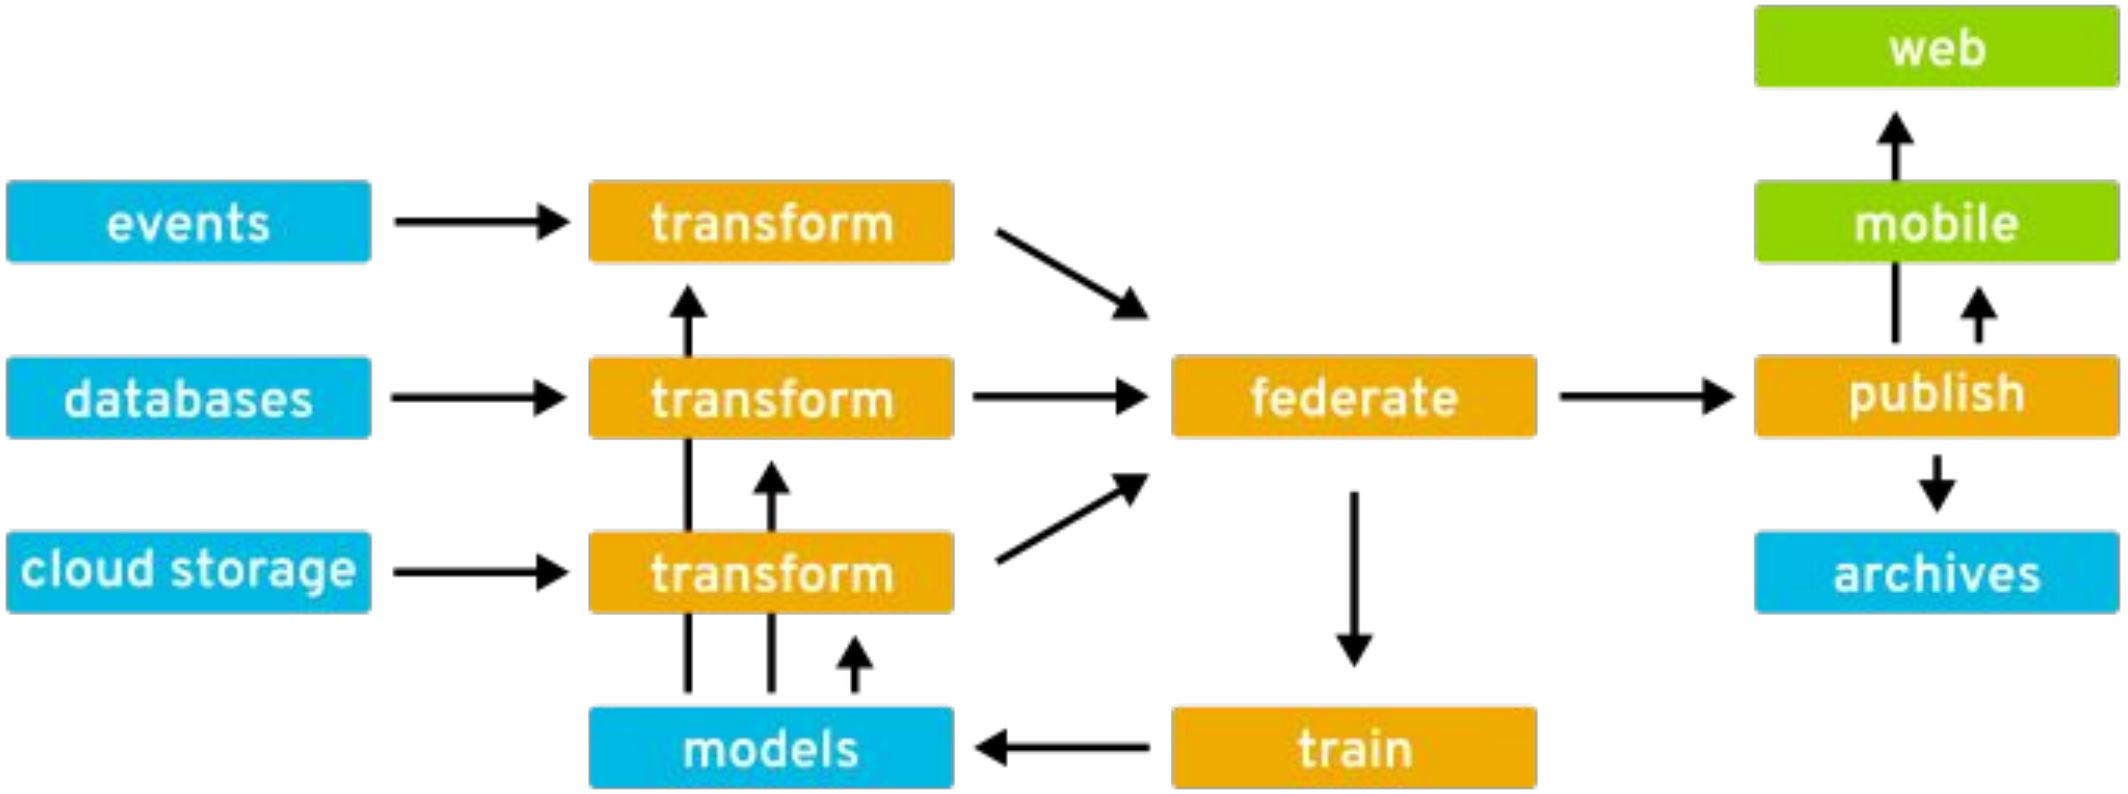

Red Hat

Service Based Components

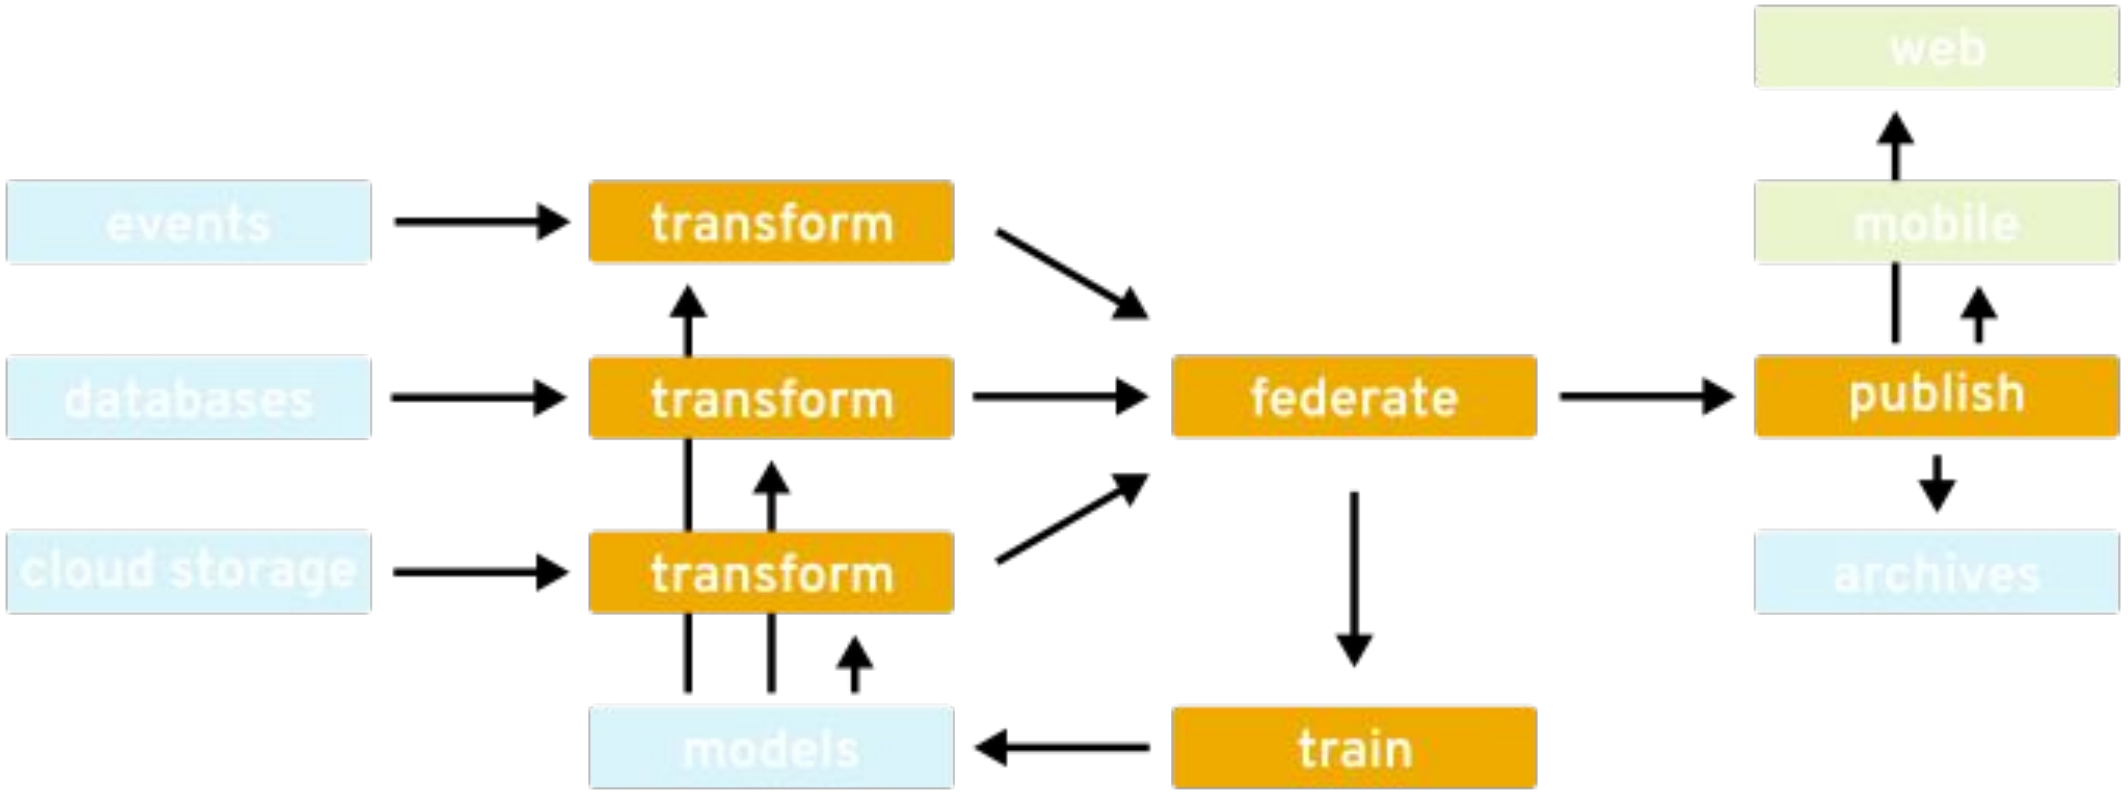

Red Hat

Intelligent Application Lifecycle

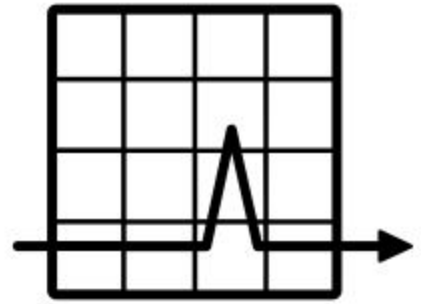

Case Study
Value at Risk

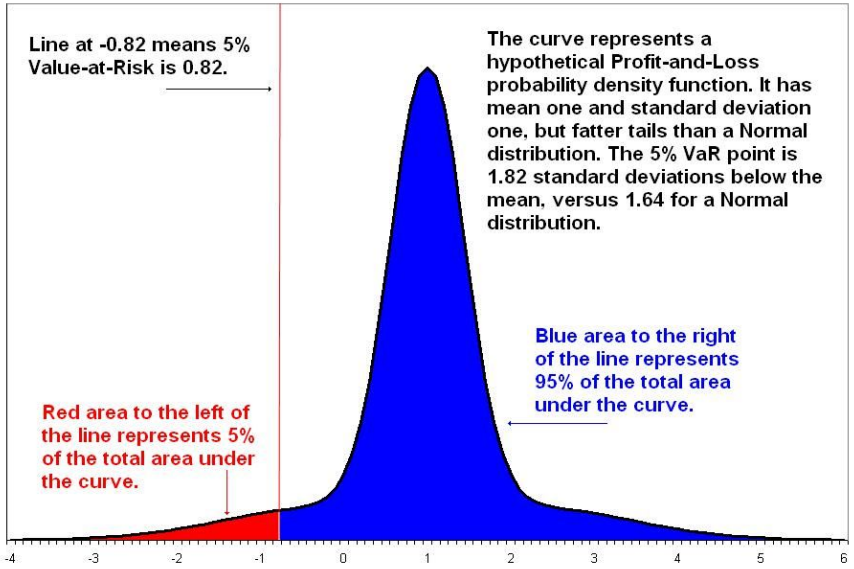

Red Hat

Going Cloud Native

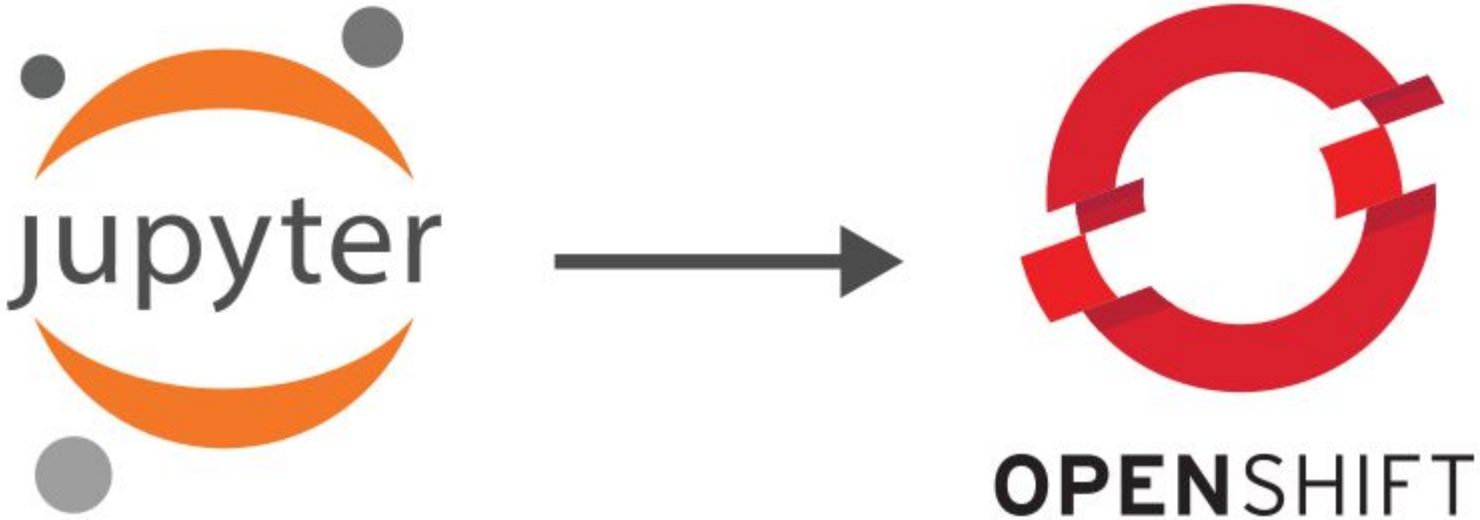

What is Cloud Native?

- Containerized
- Dynamically orchestrated
- Microservice architectures
- learn more at **cncf.io/about/faq**

What will your application do?

Red Hat

Storyboard your architecture

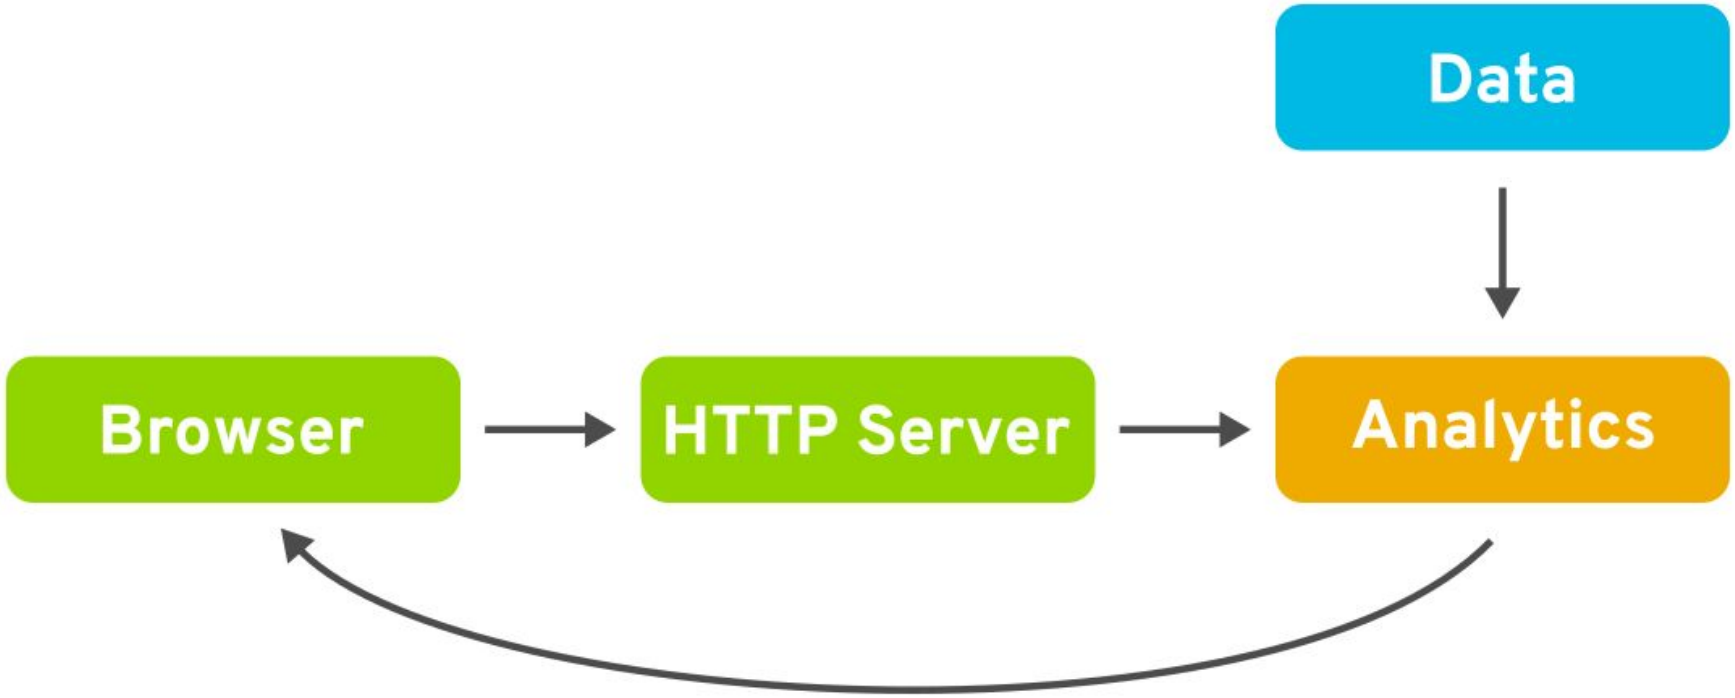

Choose wisely

Red Hat

VAR Demo

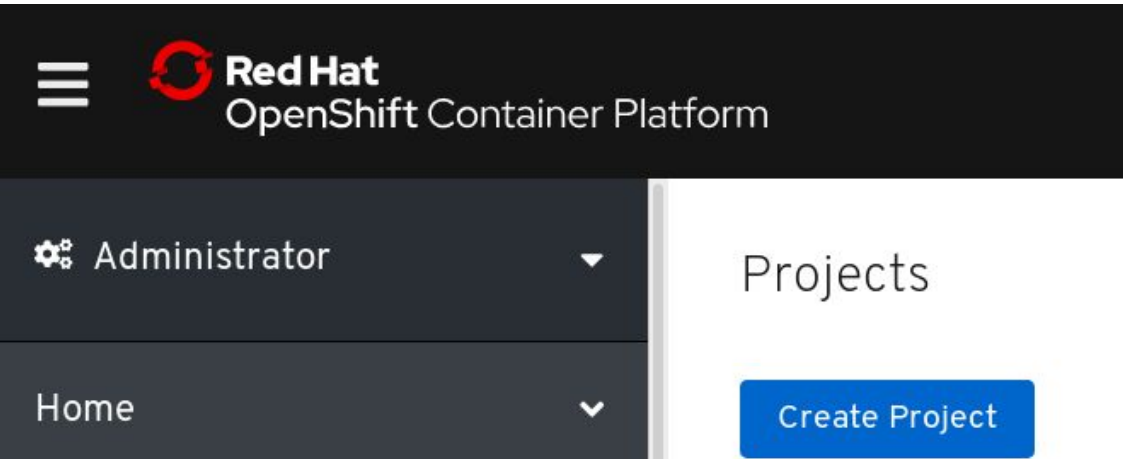

Red Hat

General architecture

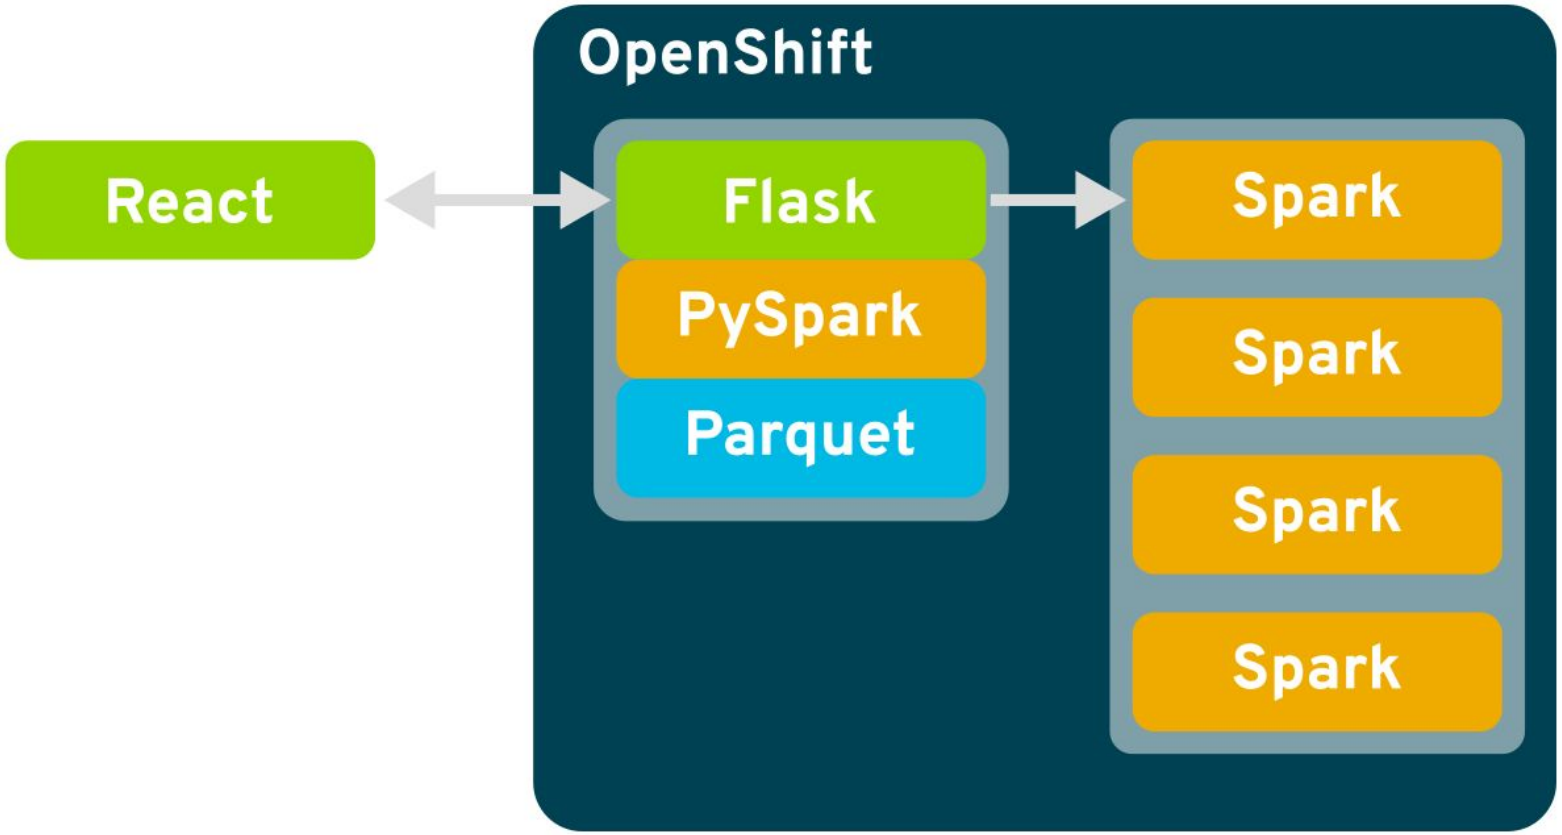

How was it built?

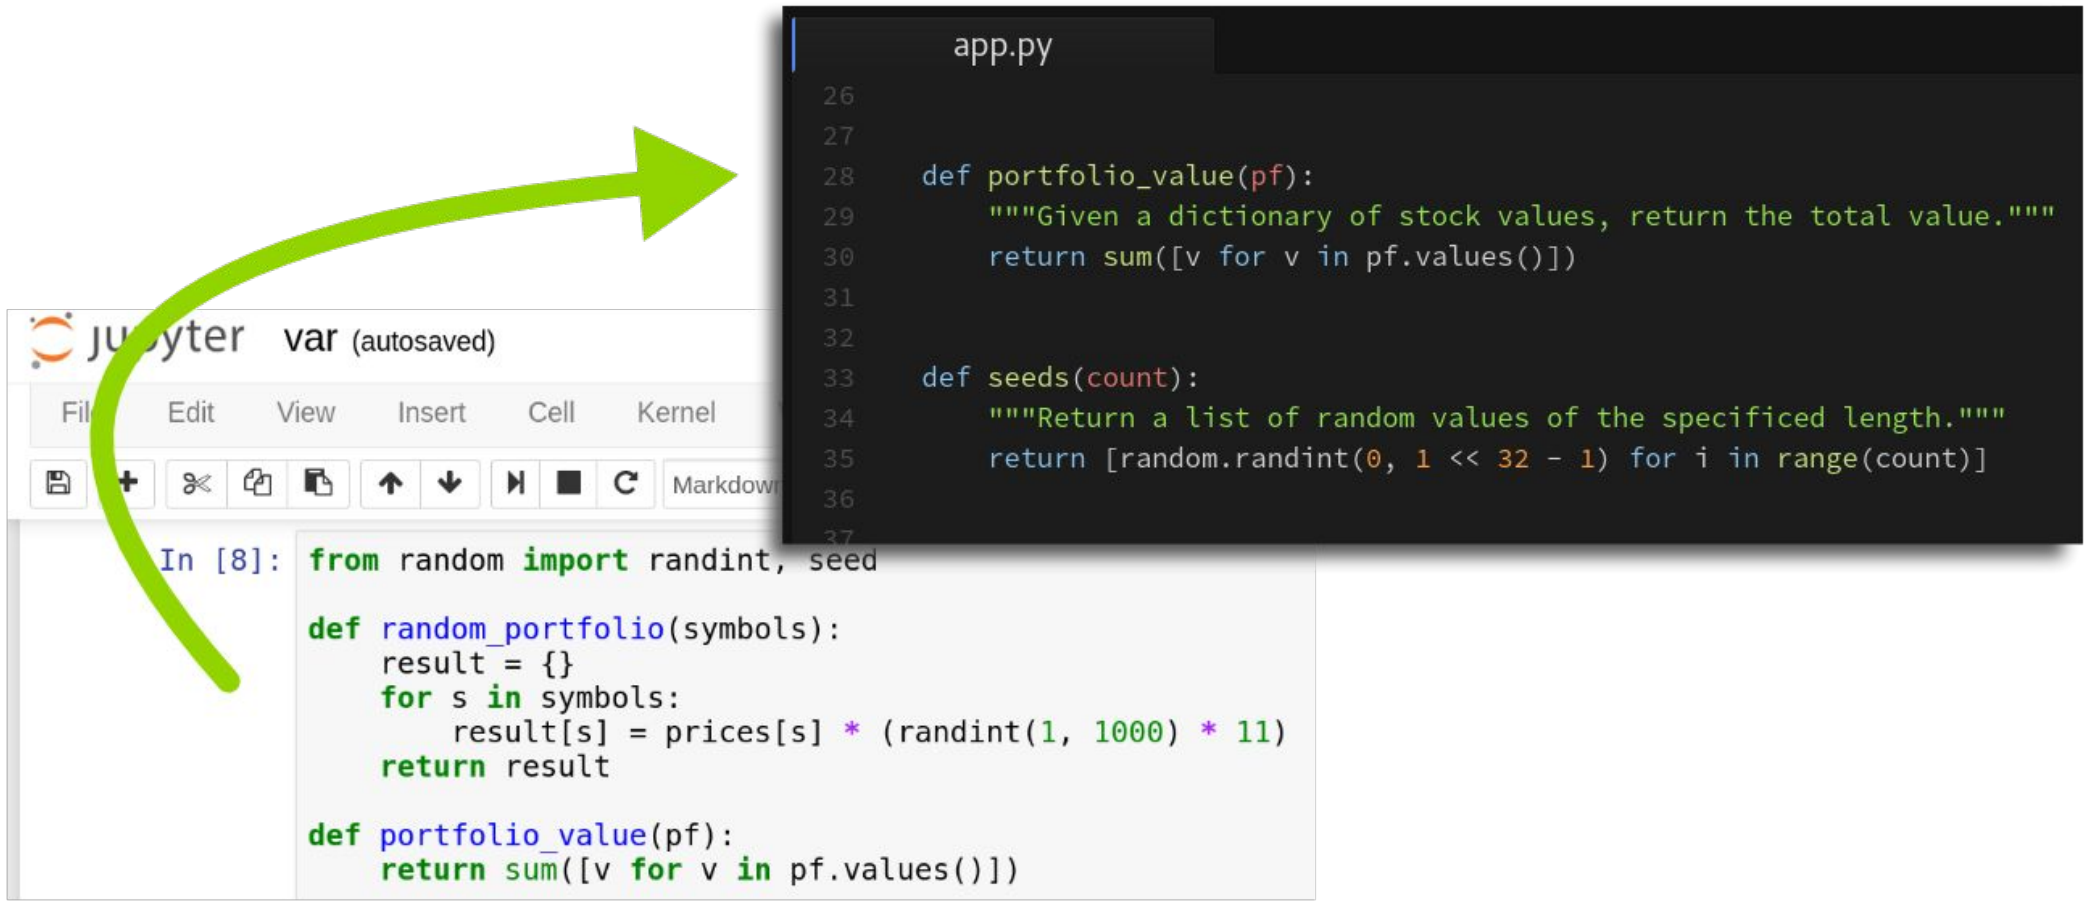

Red Hat

Moving out of alpha

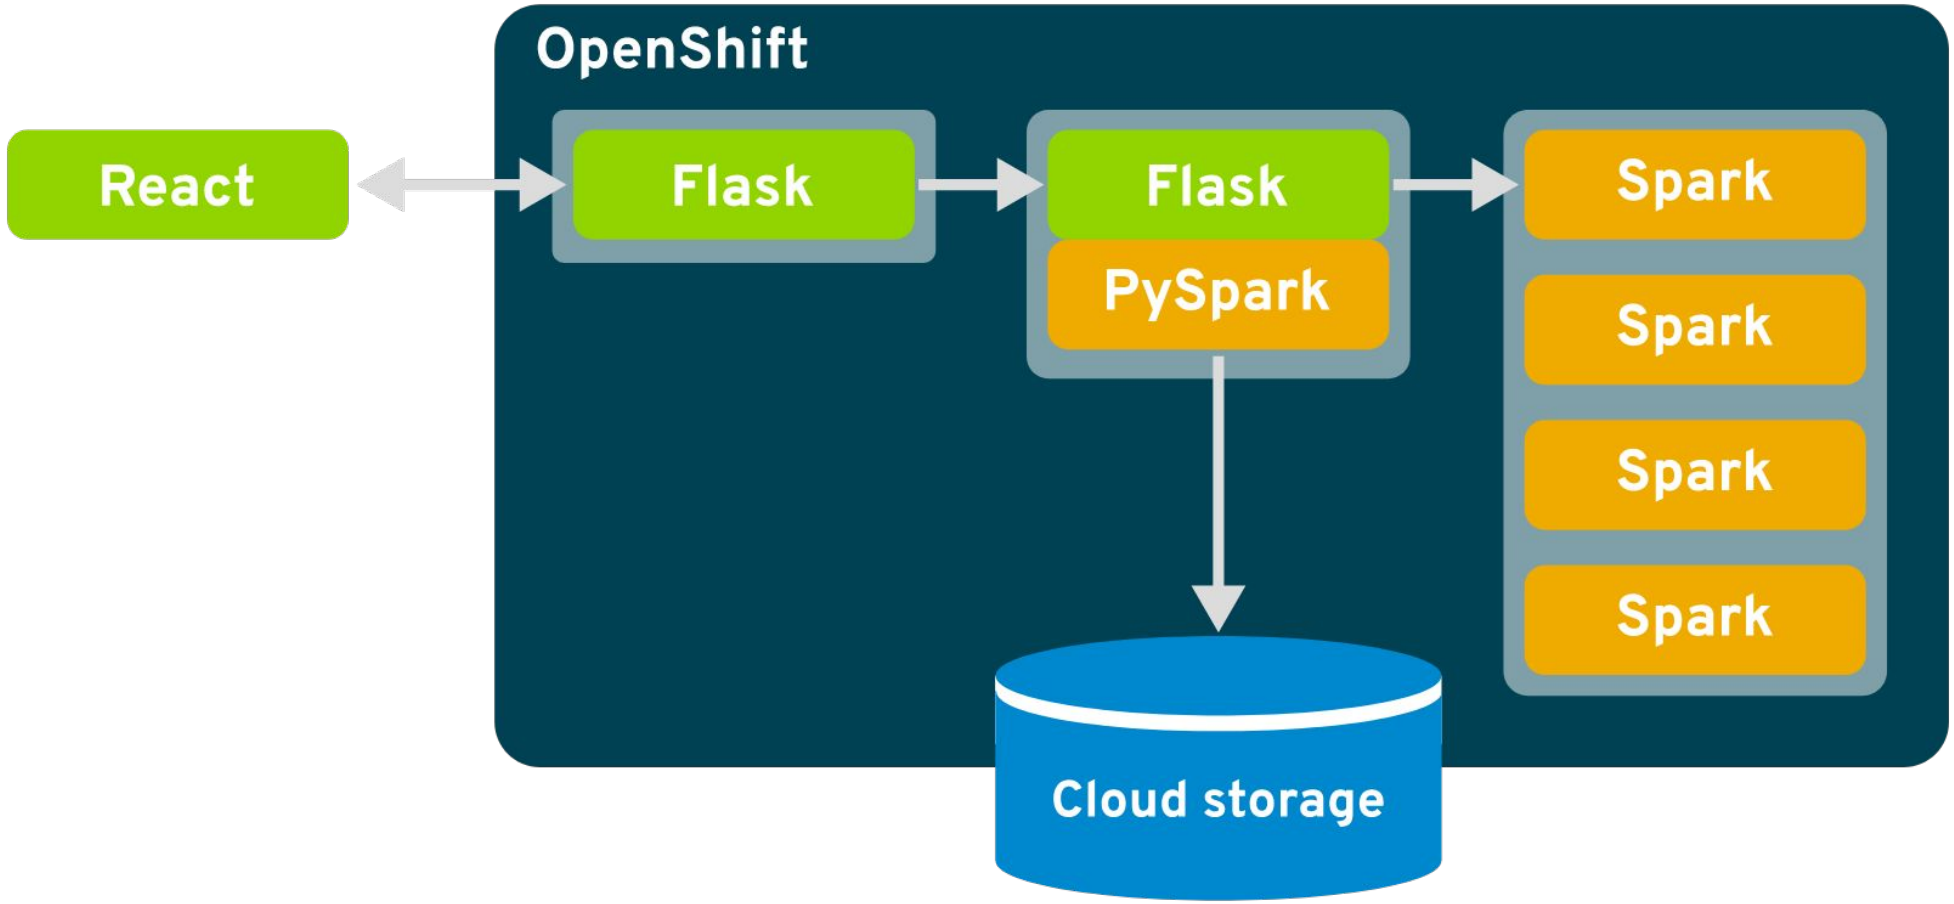

Layer in DevOps

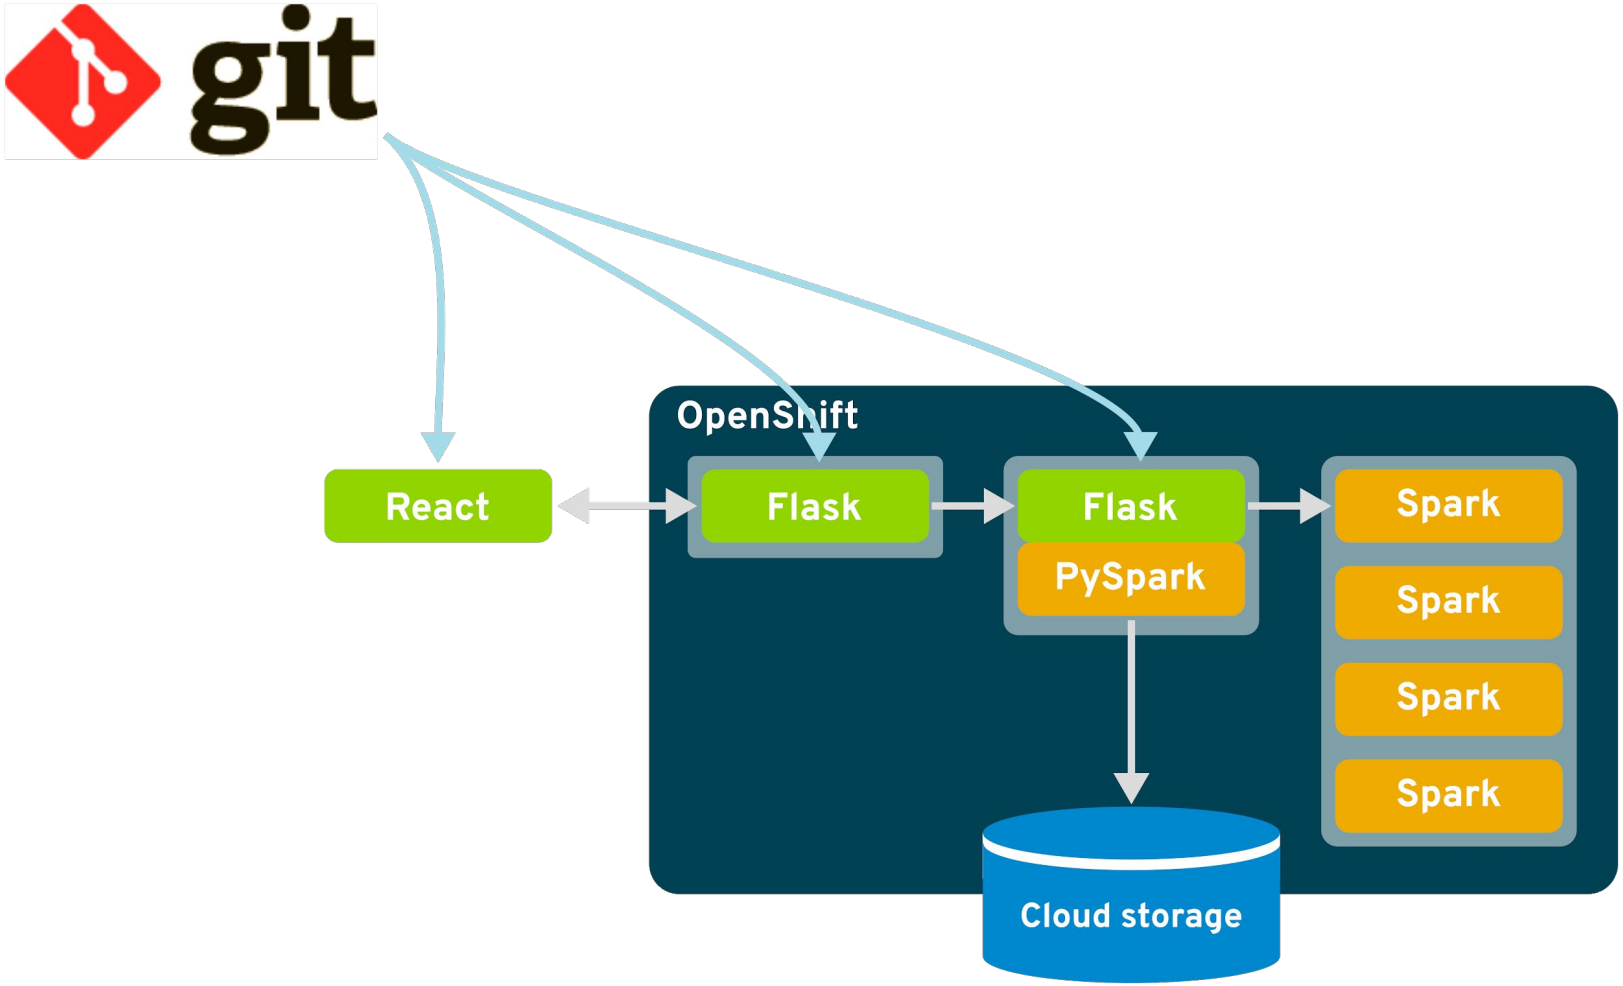

Red Hat

Augment with lifecycle monitoring

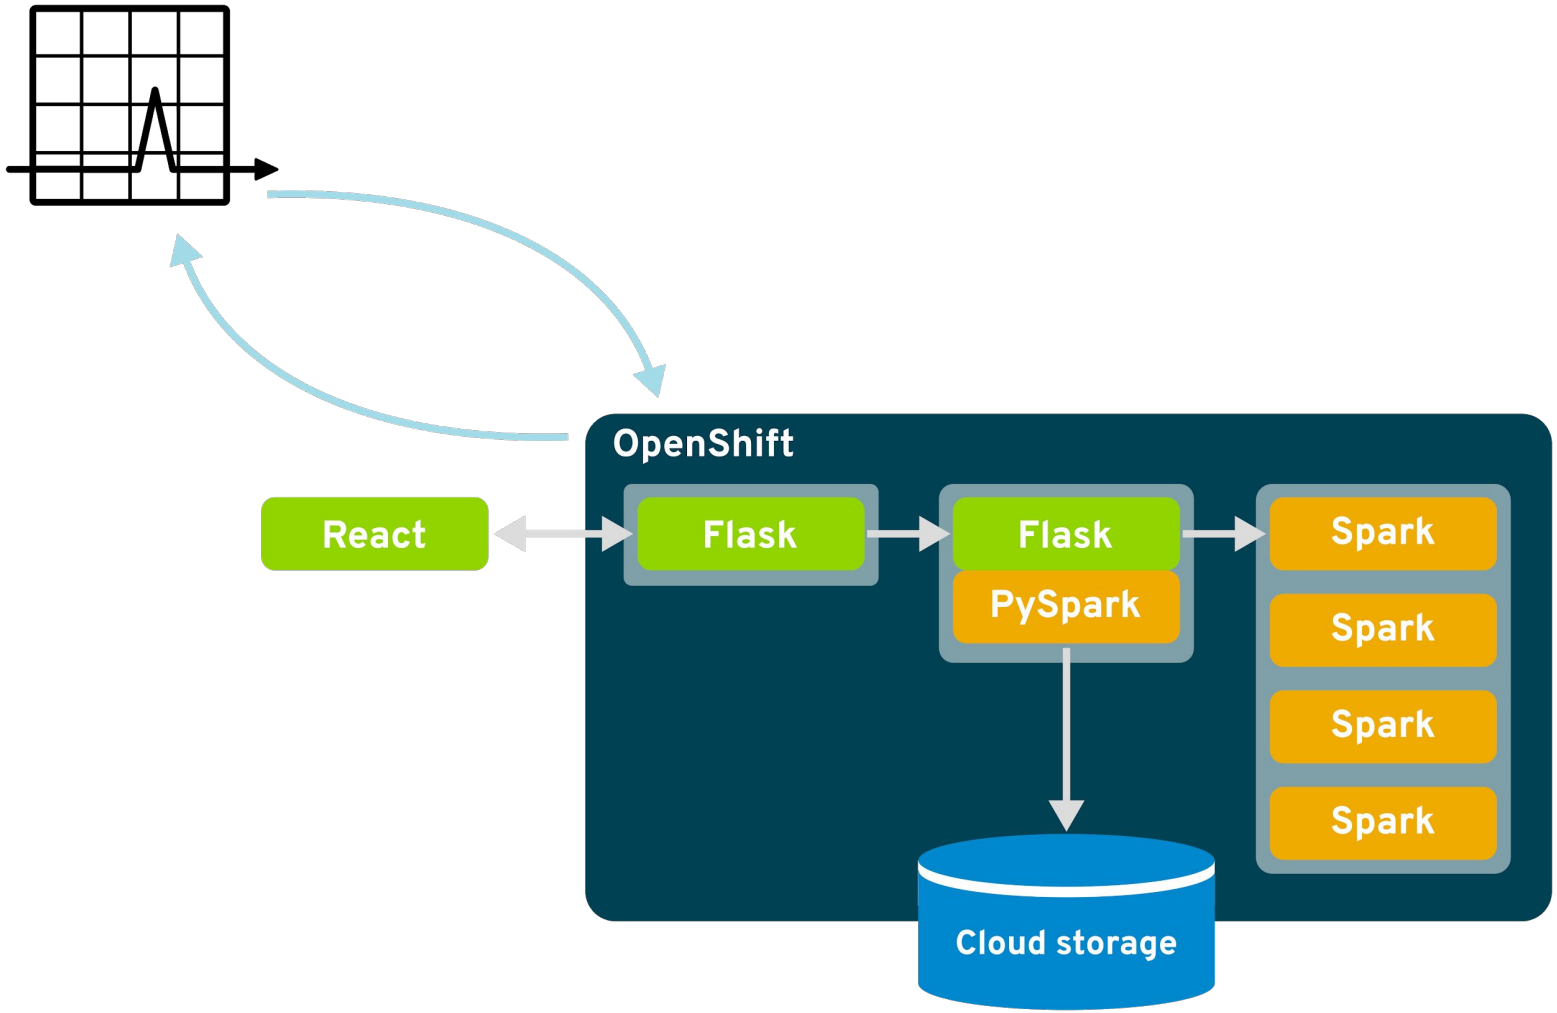

Harvest customized metrics

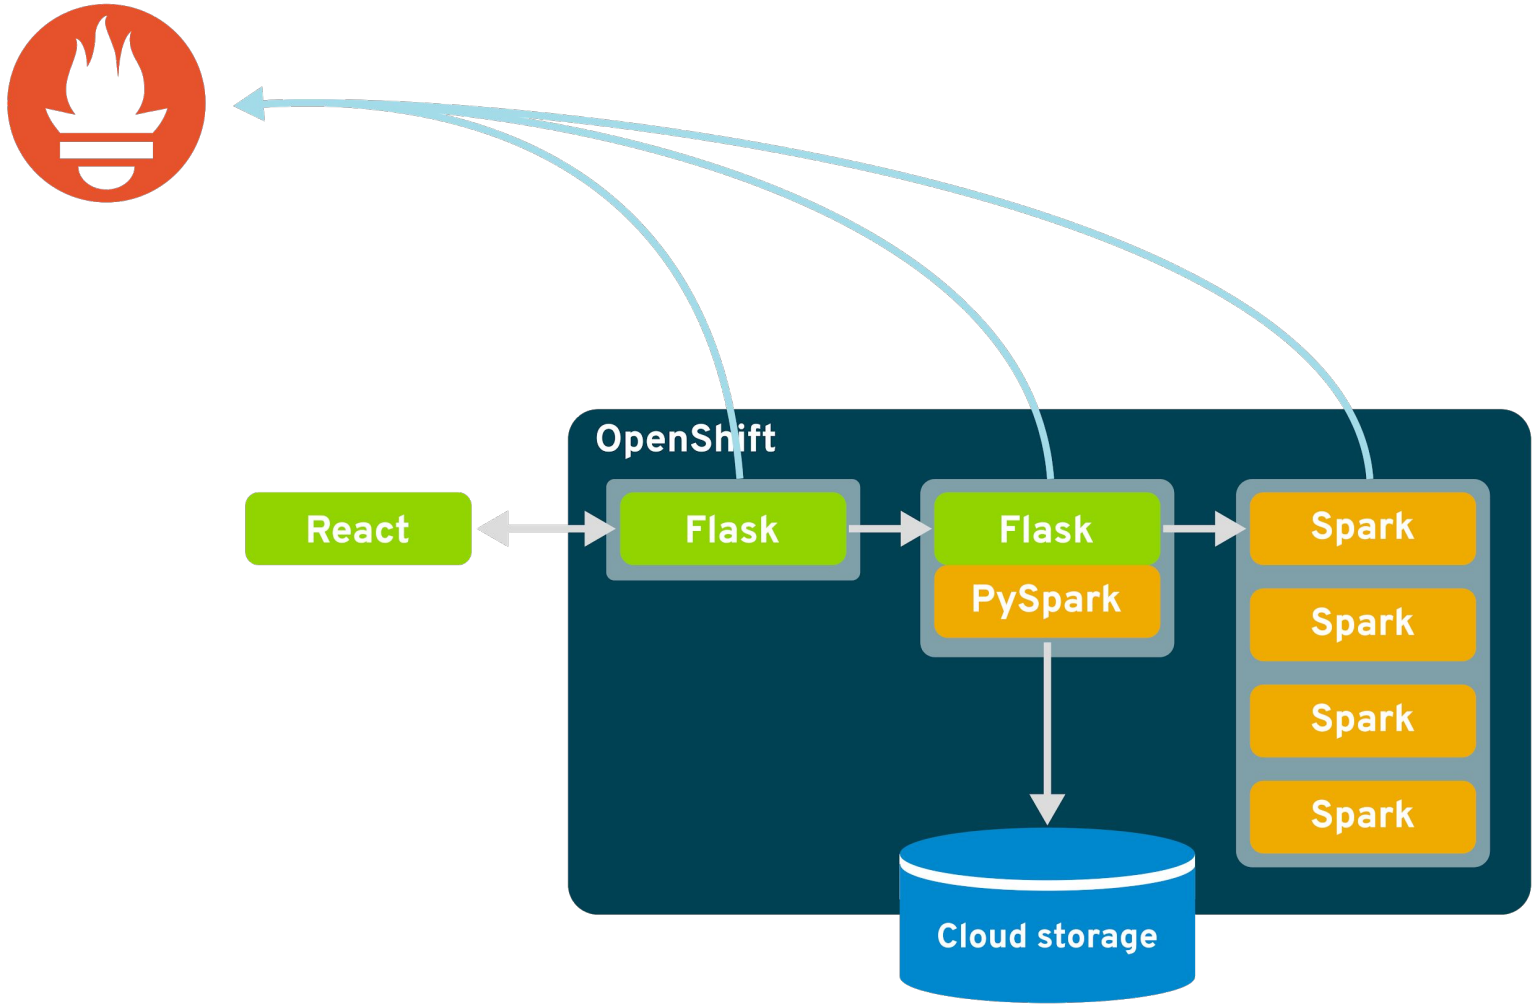

Red Hat

Lessons Learned

Red Hat

Code Comments 3.0

Value-at-risk calculations

The basic idea behind the value-at-risk calculation is that we're going to look at the historical returns of a portfolio of securities and run many simulations to determine the range of returns we can expect from these. We can then predict, over a given time horizon, what our expected loss is at a given probability, e.g., we might say that there is less than a 10% chance that the portfolio will lose more than $1,000,000.

Note that this is a didactic example and consequently makes some simplifying assumptions about the composition of the portfolio (i.e., only long positions in common stocks, so no options, dividends, or short selling) and the behavior of the market (i.e., day-to-day return percentages are normally-distributed and independent). Do not use this code to guide actual investment decisions!

Basic setup

Here we import the `pyspark` module and set up a `SparkSession`.

In [1]:
```
import pyspark
from pyspark.context import SparkContext
from pyspark.sql import SparkSession, SQLContext

spark = SparkSession.builder.master("local[*]").getOrCreate()
```

Templated Repeatability

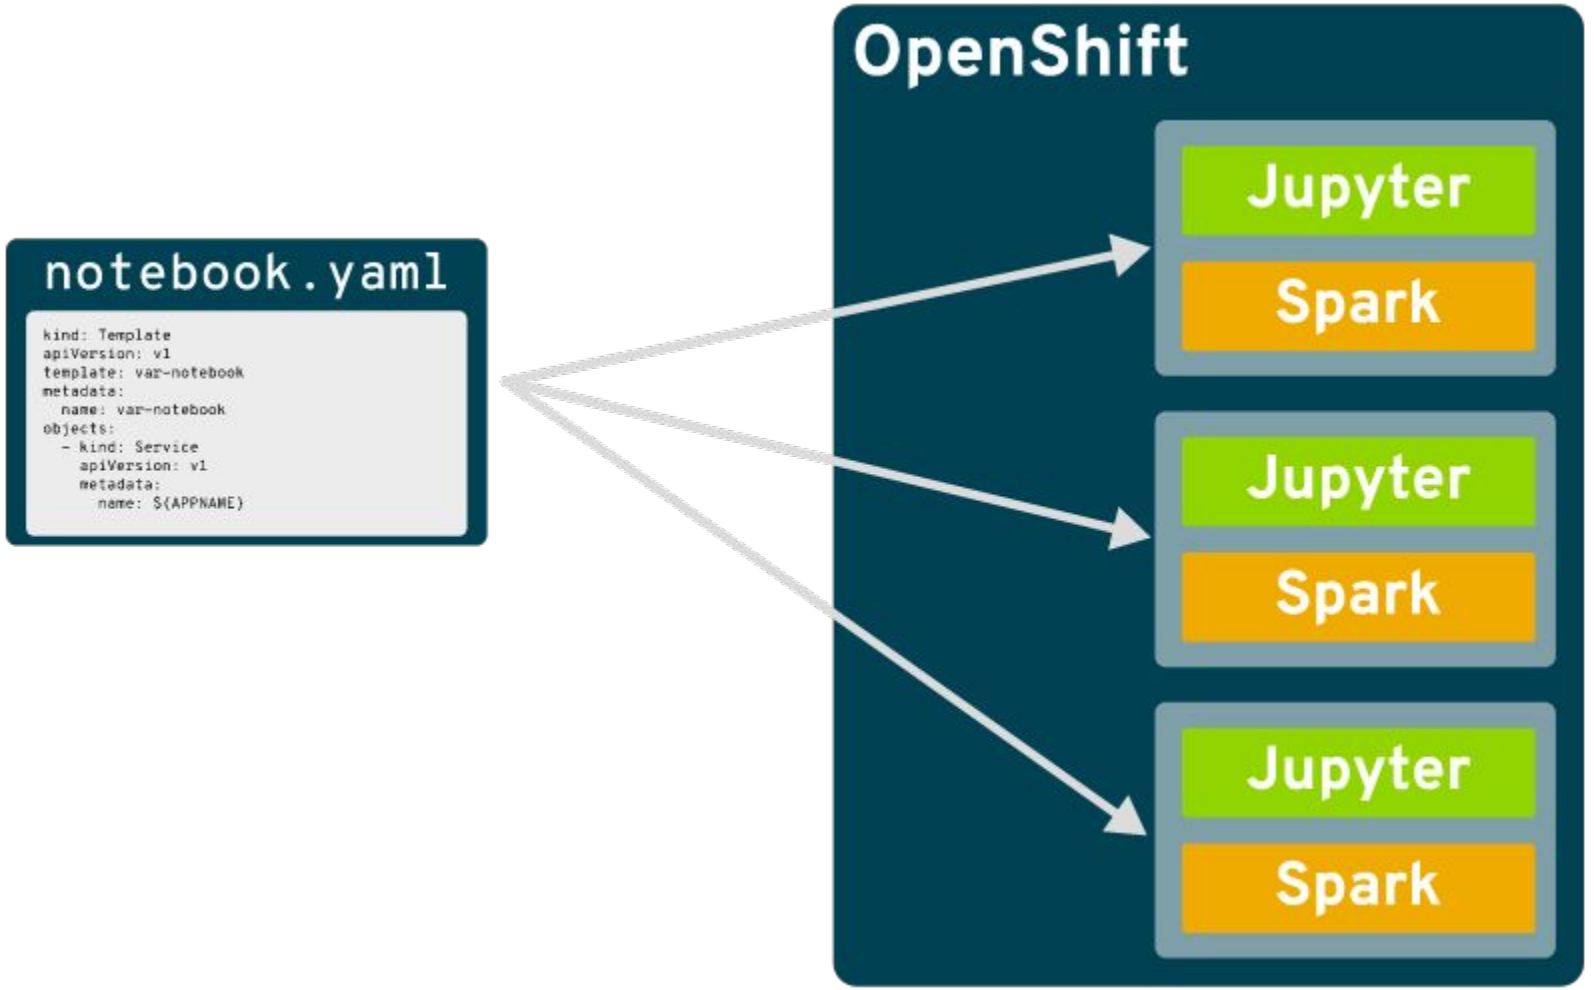

Red Hat

Your greatest tool on the journey

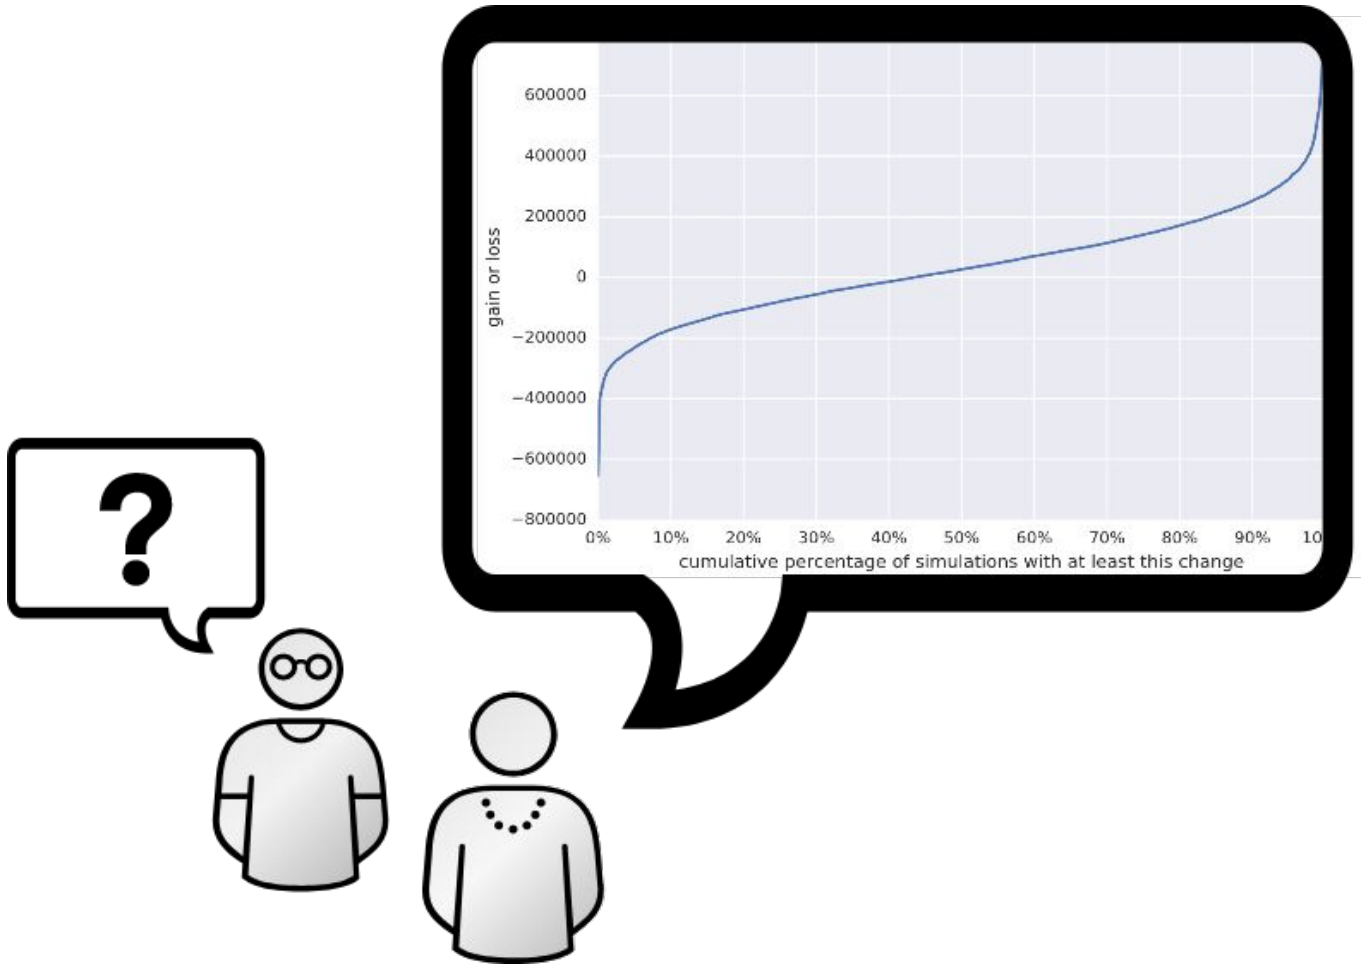

Red Hat

Second greatest tool?

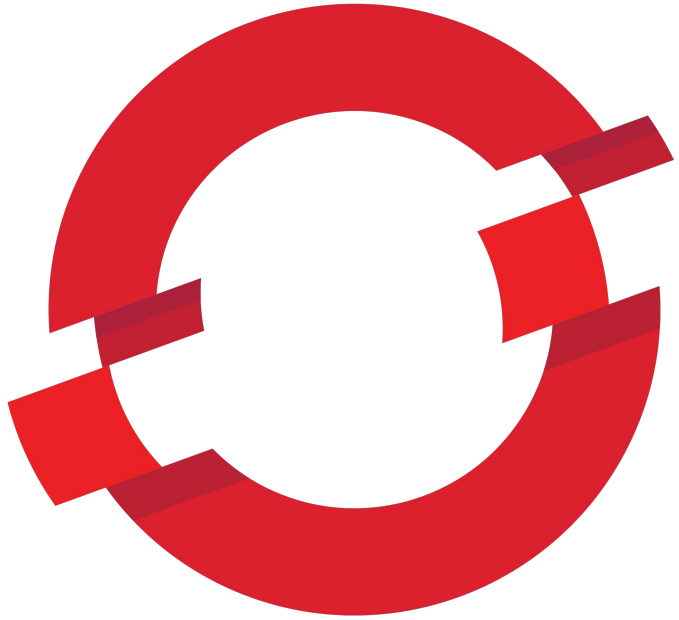

OPENSHIFT

What Next?

AI/ML on OpenShift

www.openshift.com/learn/topics/ai-ml

Open Data Hub

opendatahub.io

Red Hat

Radanalytics

radanalytics.io

Thank you

Keep in touch

elmiko@redhat.com

@elmiko@mastodon.technology

github.com/elmiko/rhug-artifacts

Red Hat